

The Docker Way

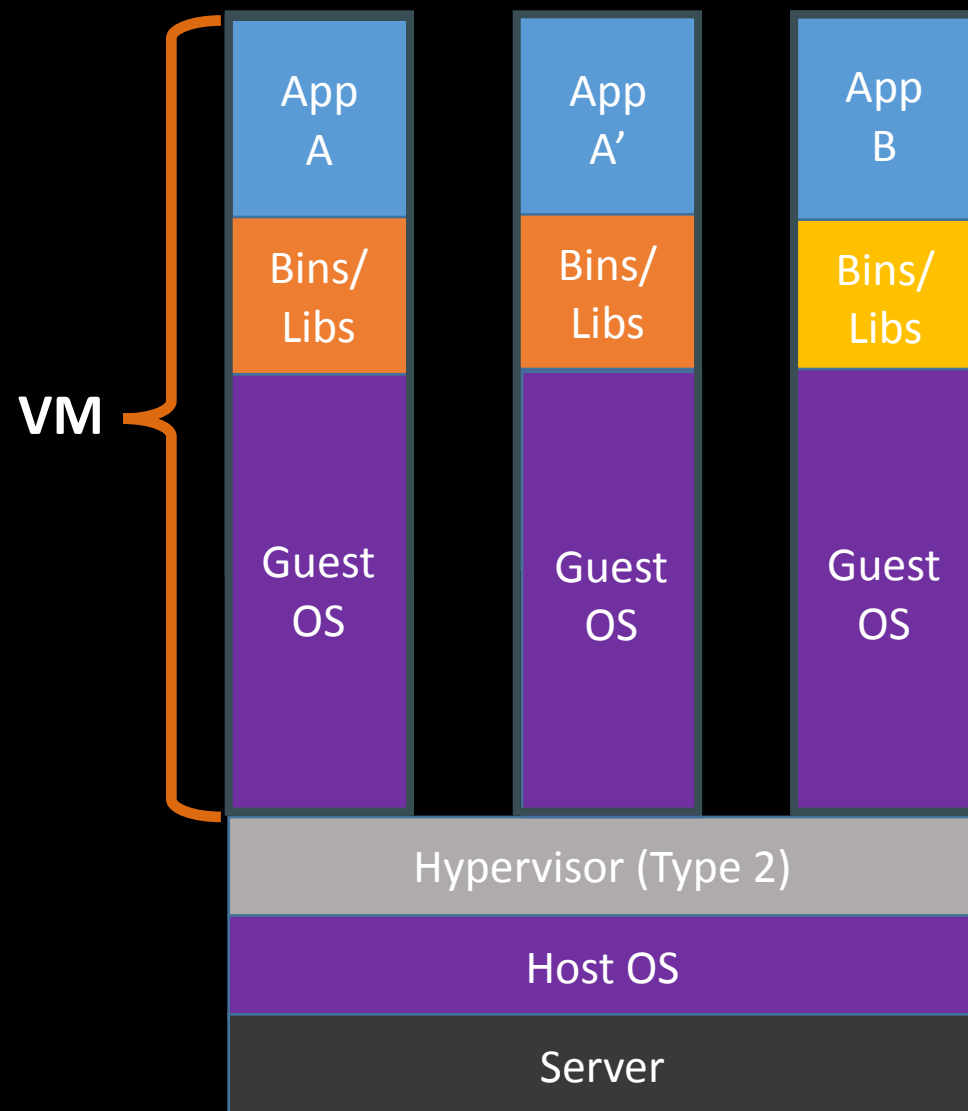
xuli@qiniu.com

Agenda

- Docker: What, Why and How
- Story: Qiniu meet Docker
- Q & A

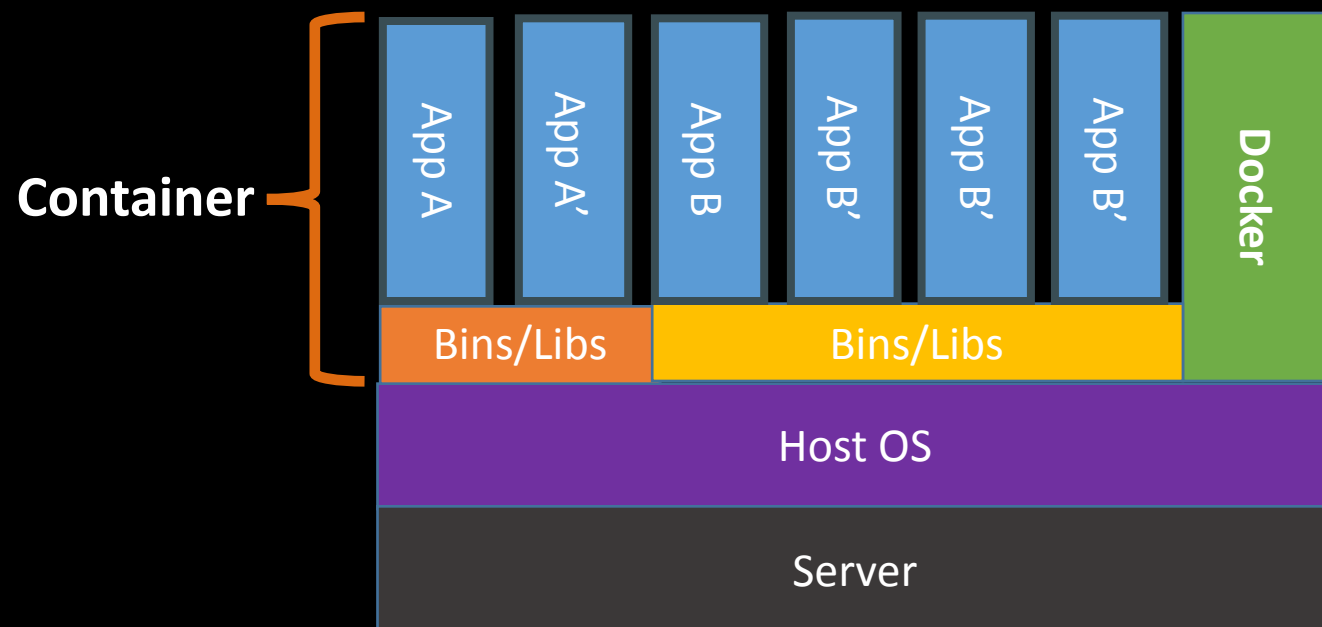
Container do Virtualization

Container vs. VM



Containers are isolated, but share OS and, where appropriate, bins/libraries

...result is significantly faster deployment, much less overhead, easier migration, faster restart



- The same see and shared for each Linux-process:
 - Linux Kernel
 - File System
 - Network System
 - PID, UID, IPC, etc..
 - Memory, Disk, CPU, etc..

Virtualization

- Isolated:
 - FS, net, pid, uid, uts, mnt, ipc namespace
- Constrained:
 - Memory, CPU, Network IO, Disk IO, Disk Space

The process groups that meet the
above restrictions are called
**"lightweight virtual machine" or
Container**

Why Docker ?

The Challenge

Multiplicity of Stacks



Static website

nginx 1.5 + modsecurity + openssl +
bootstrap 3



Background workers

Python 3.0 + celery + pyredis + libcurl + ffmpeg +
libopencv + nodejs + phantomjs



User DB

postgresql + pgv8 + v8



Queue

Redis + redis-sentinel



Analytics DB

hadoop + hive + thrift + OpenJDK



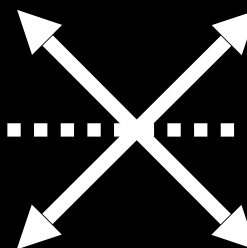
Web frontend

Ruby + Rails + sass + Unicorn



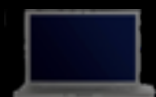
API endpoint

Python 2.7 + Flask + pyredis + celery + pycopg + postgresql-
client



Do services and apps
interact appropriately?

Multiplicity of
hardware
environments



Development VM



QA server

Customer Data Center



Public Cloud

Disaster recovery

Production Servers



Production Cluster
















Contributor's laptop



Can I migrate smoothly
and quickly?

N x N compatibility nightmare

	Static website	?	?	?	?	?	?	?
	Web frontend	?	?	?	?	?	?	?
	Background workers	?	?	?	?	?	?	?
	User DB	?	?	?	?	?	?	?
	Analytics DB	?	?	?	?	?	?	?
	Queue	?	?	?	?	?	?	?
		Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
								

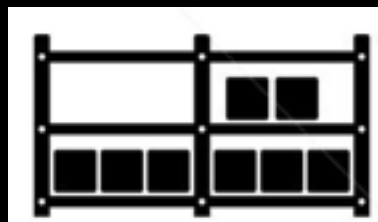
An Inspiration.. and history: Cargo Transport Pre-1960

Multiplicity of Goods



Do I worry about how
goods interact (e.g.
coffee beans next to
spices)

Multiplicity of methods
for transporting/storing



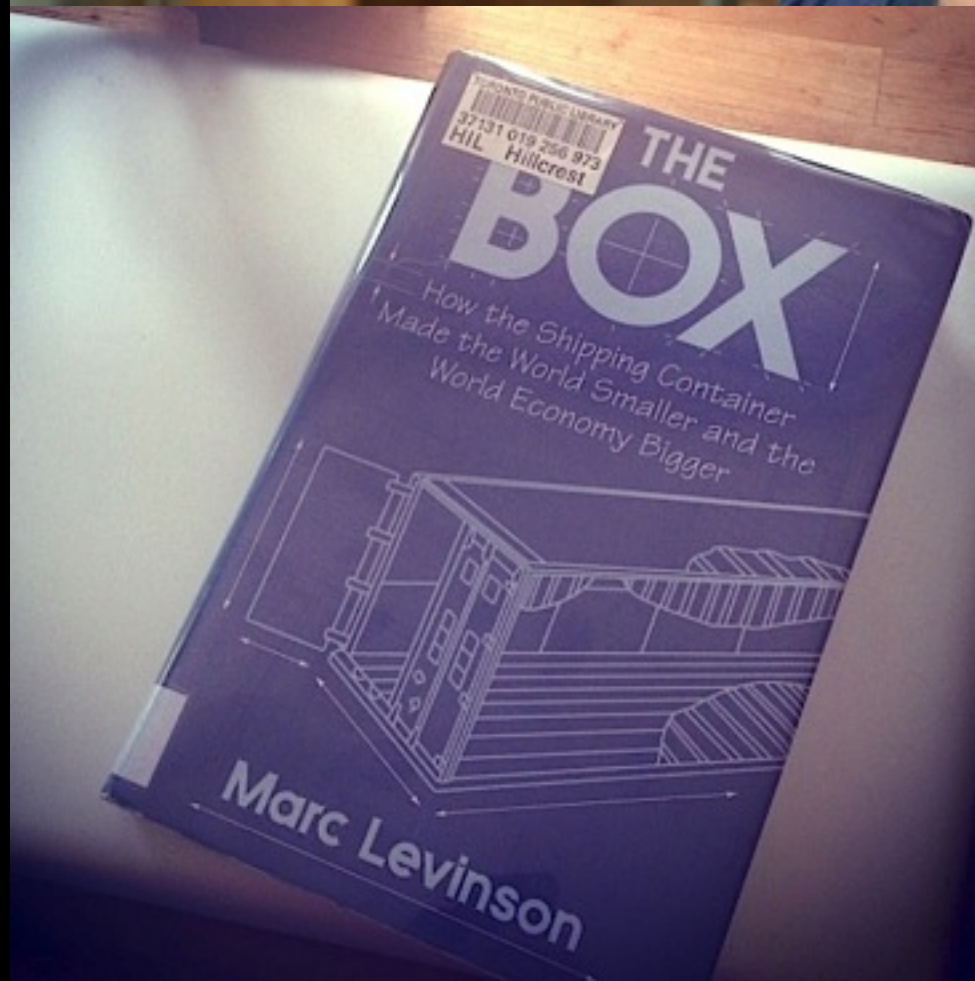
Can I transport quickly
and smoothly
(e.g. from boat to train to
truck)

Also a matrix from hell

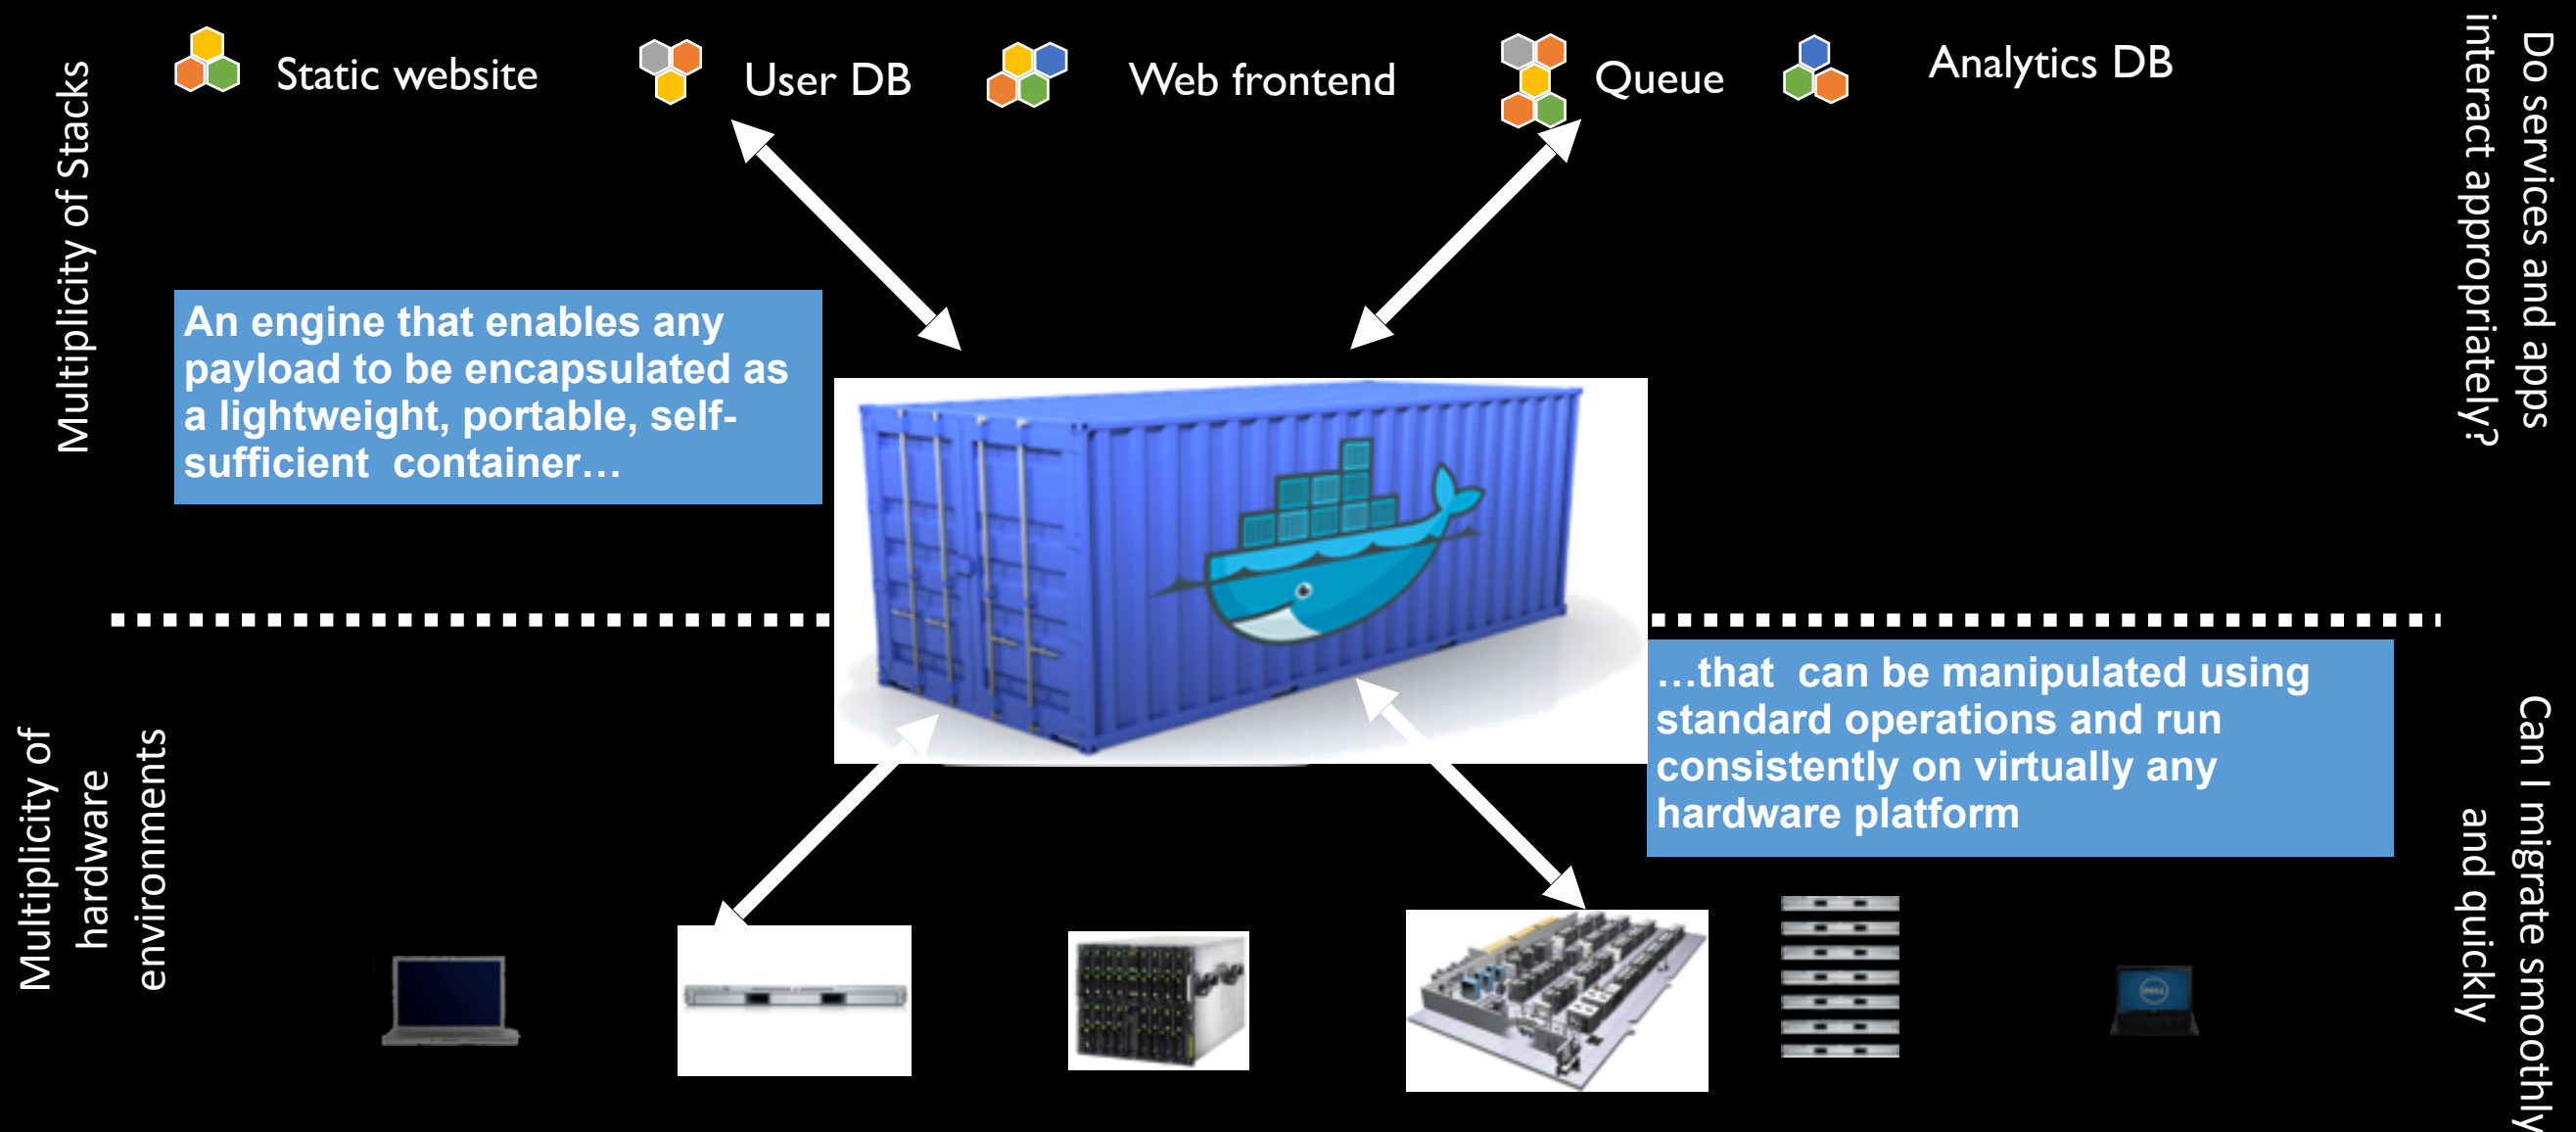
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
	?	?	?	?	?	?	?
							

Solution: Intermodal Shipping Container

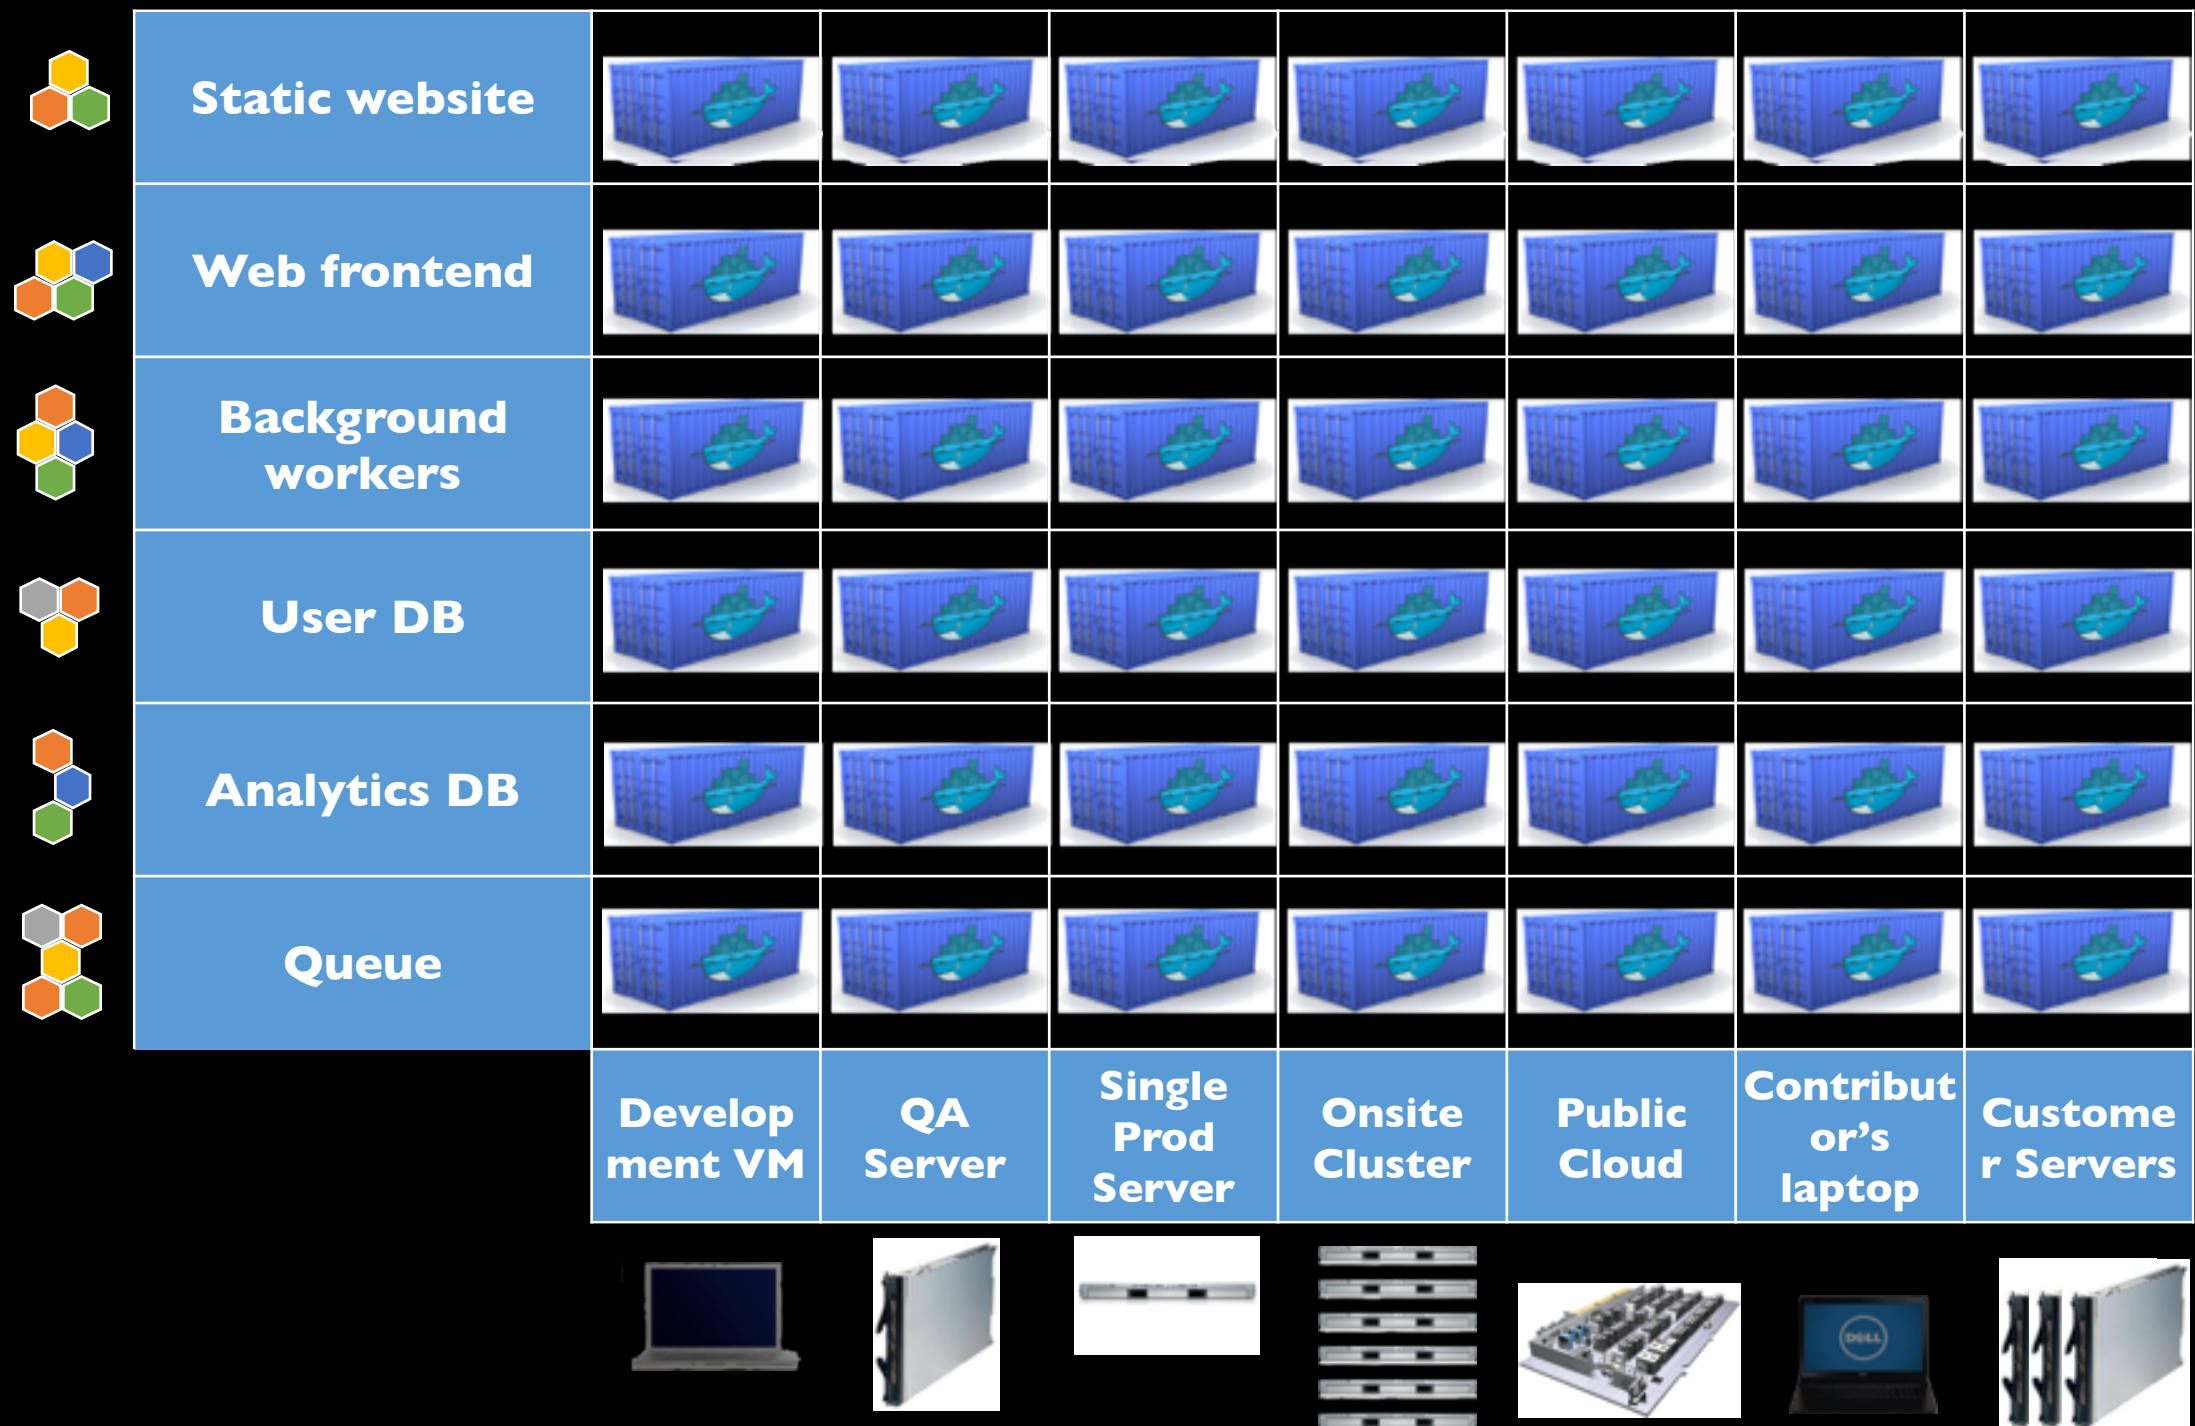




A shipping container system for code



This eliminated the N x N problem



Build, Ship, Run

“build once, run anywhere”

– *developers*

“configure once, run anything”

– *operations*

What can I use Docker for?

- Continuous Integration / Deployment
- Packaging and deploying applications
- Build your own PAAS
- Deploy applications at hyperscale

Work with Puppet or Chef

- Chef and Puppet are state management tools
- Docker images are version controlled and layered
- Small, self-contained and lightweight

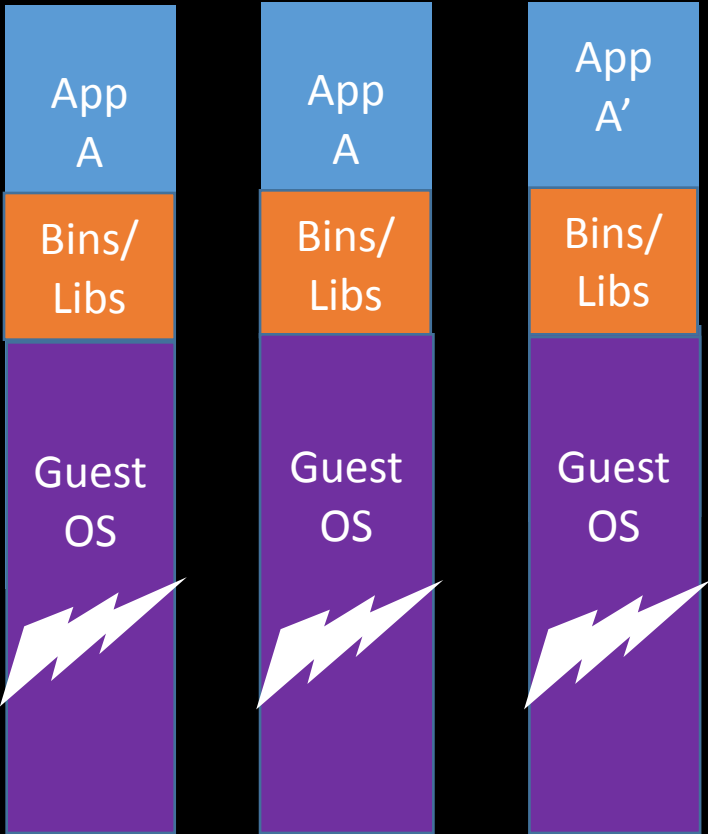
standardized,
interoperable,
automatable

Why not virtual machines?

- Speed of deployment
- Portability
- Size aka cached layering FTW
- Density & Performance
- Cost

Why are Docker containers lightweight?

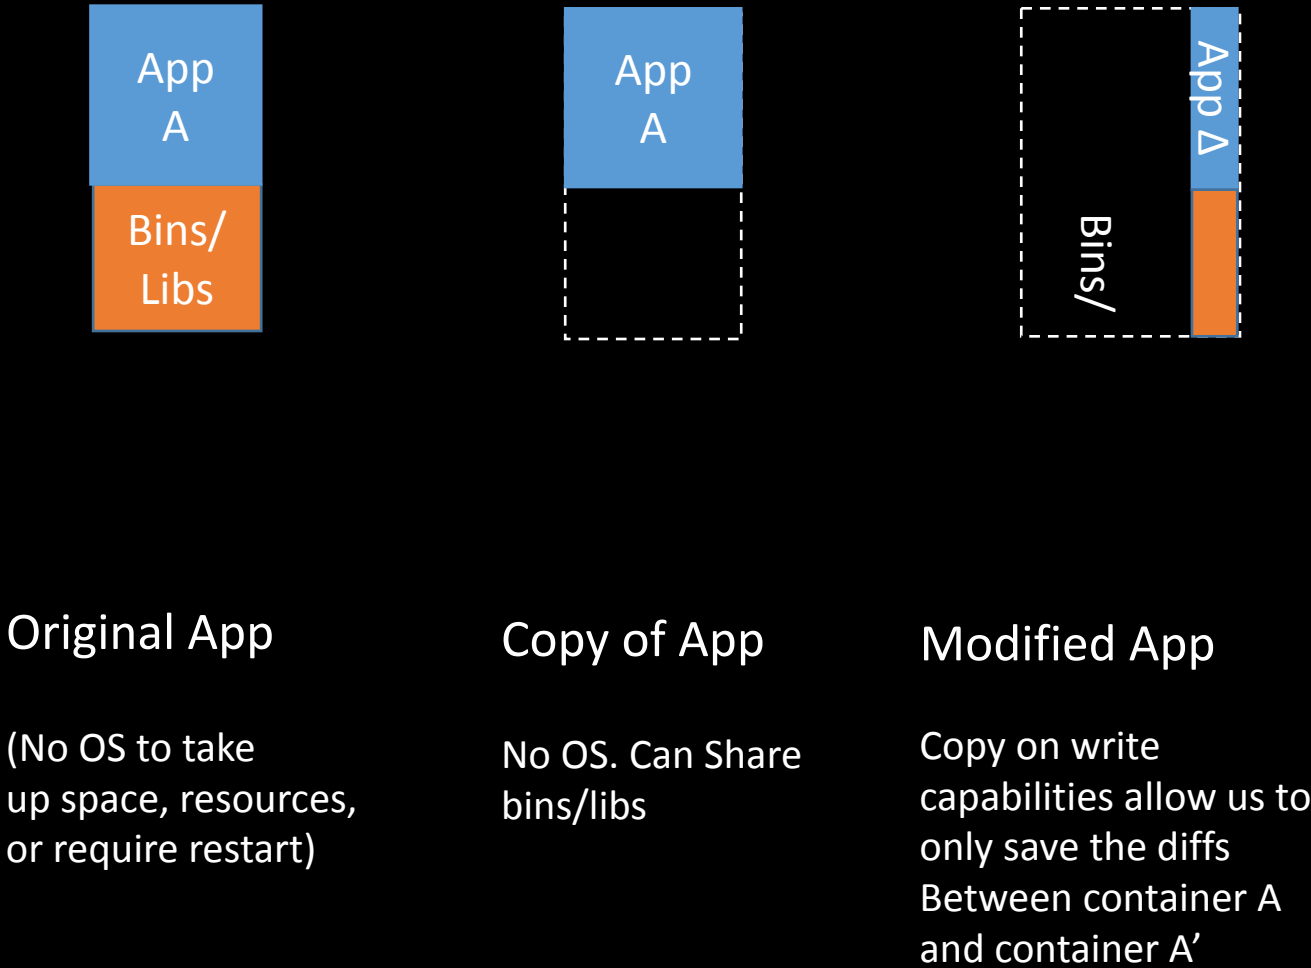
VMs



VMs

Every app, every copy of an app, and every slight modification of the app requires a new virtual server

Containers

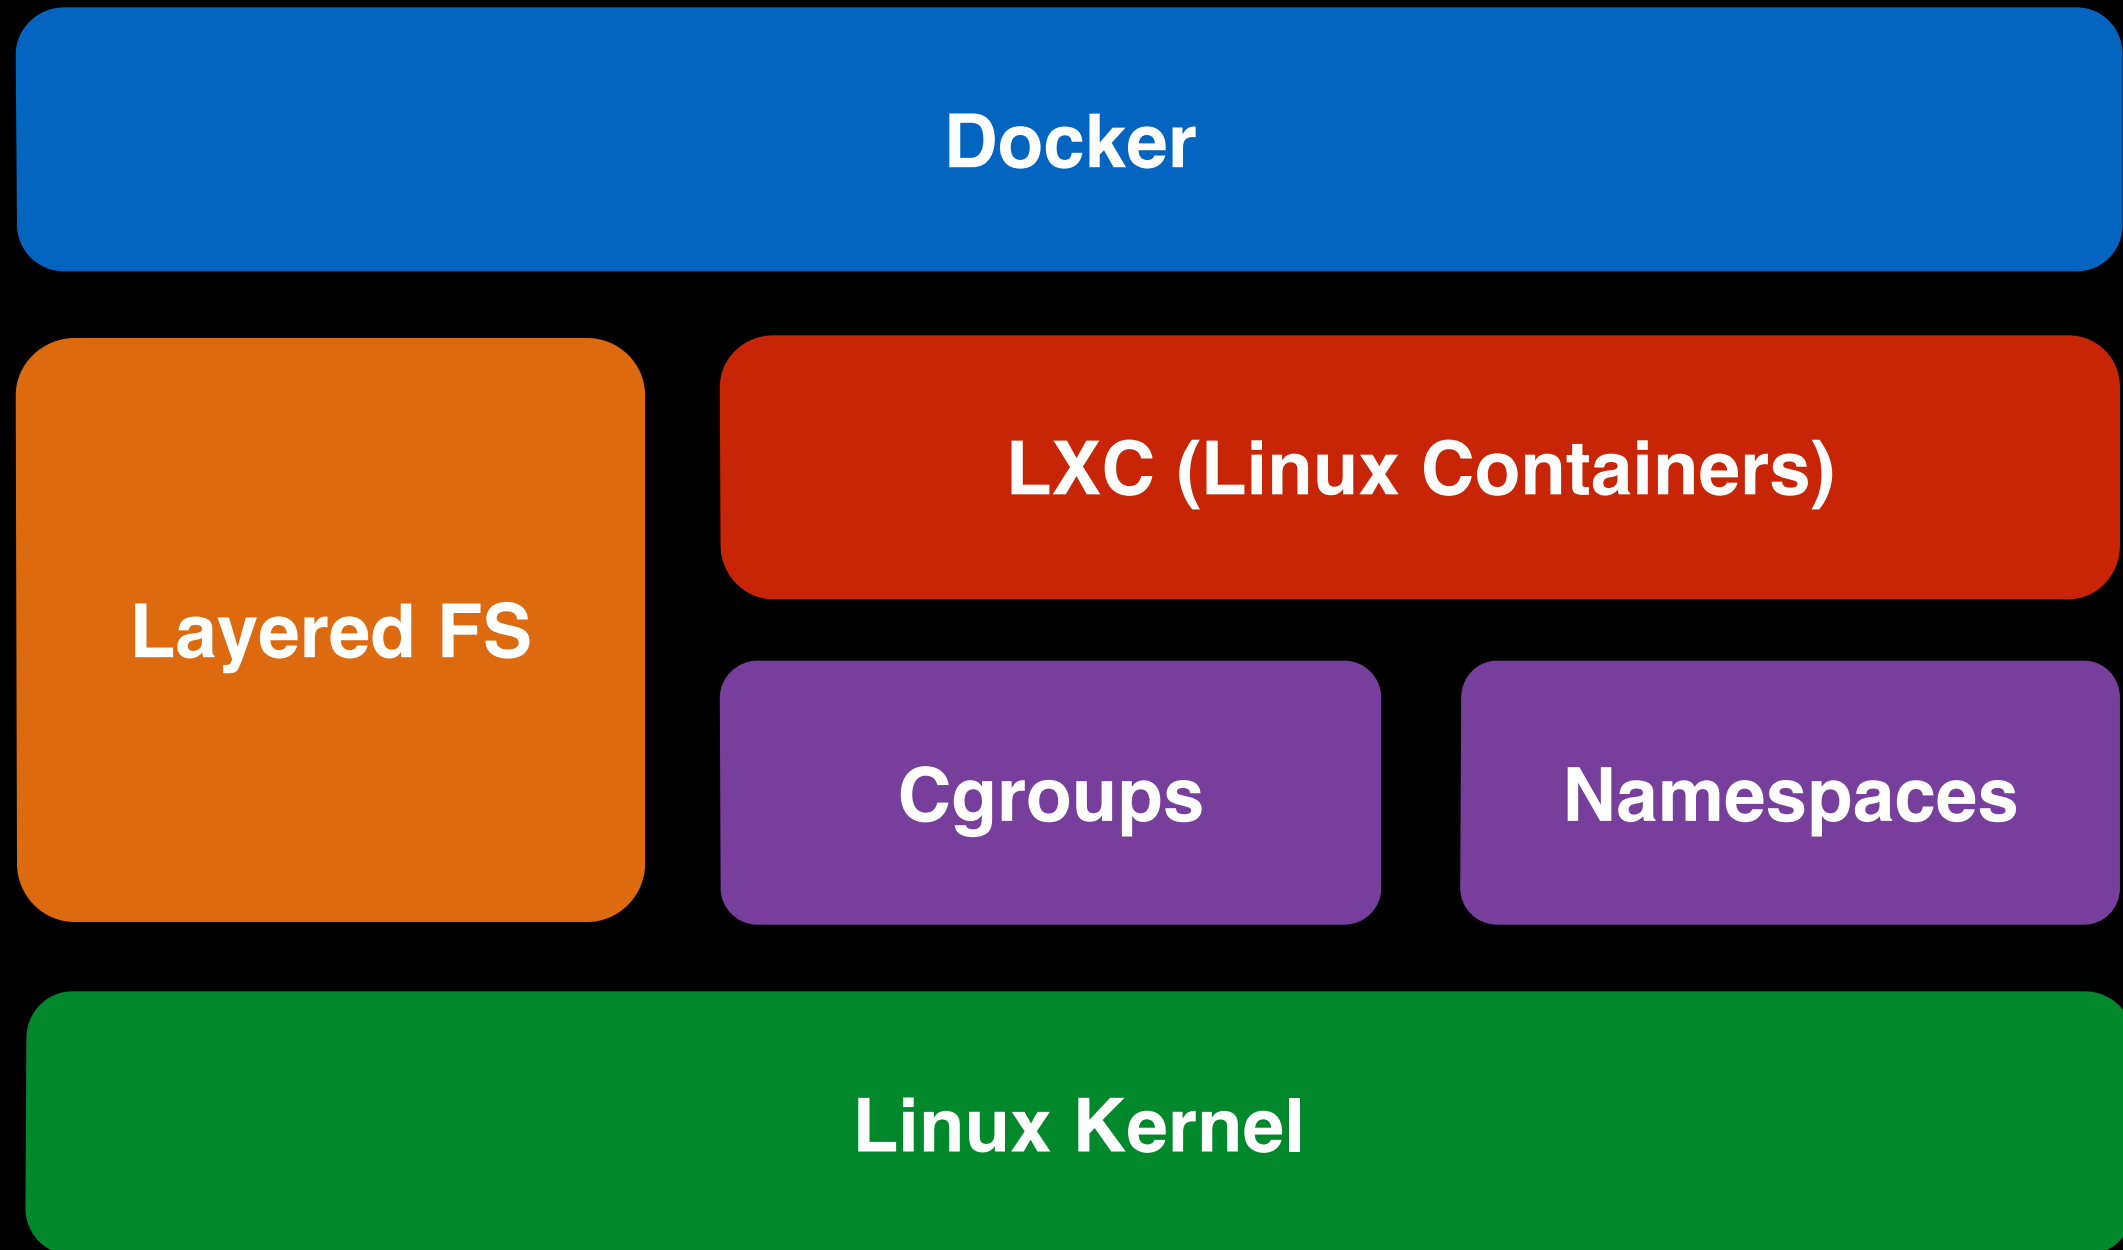


VM for IAAS,
Container for PAAS

Technology Stack

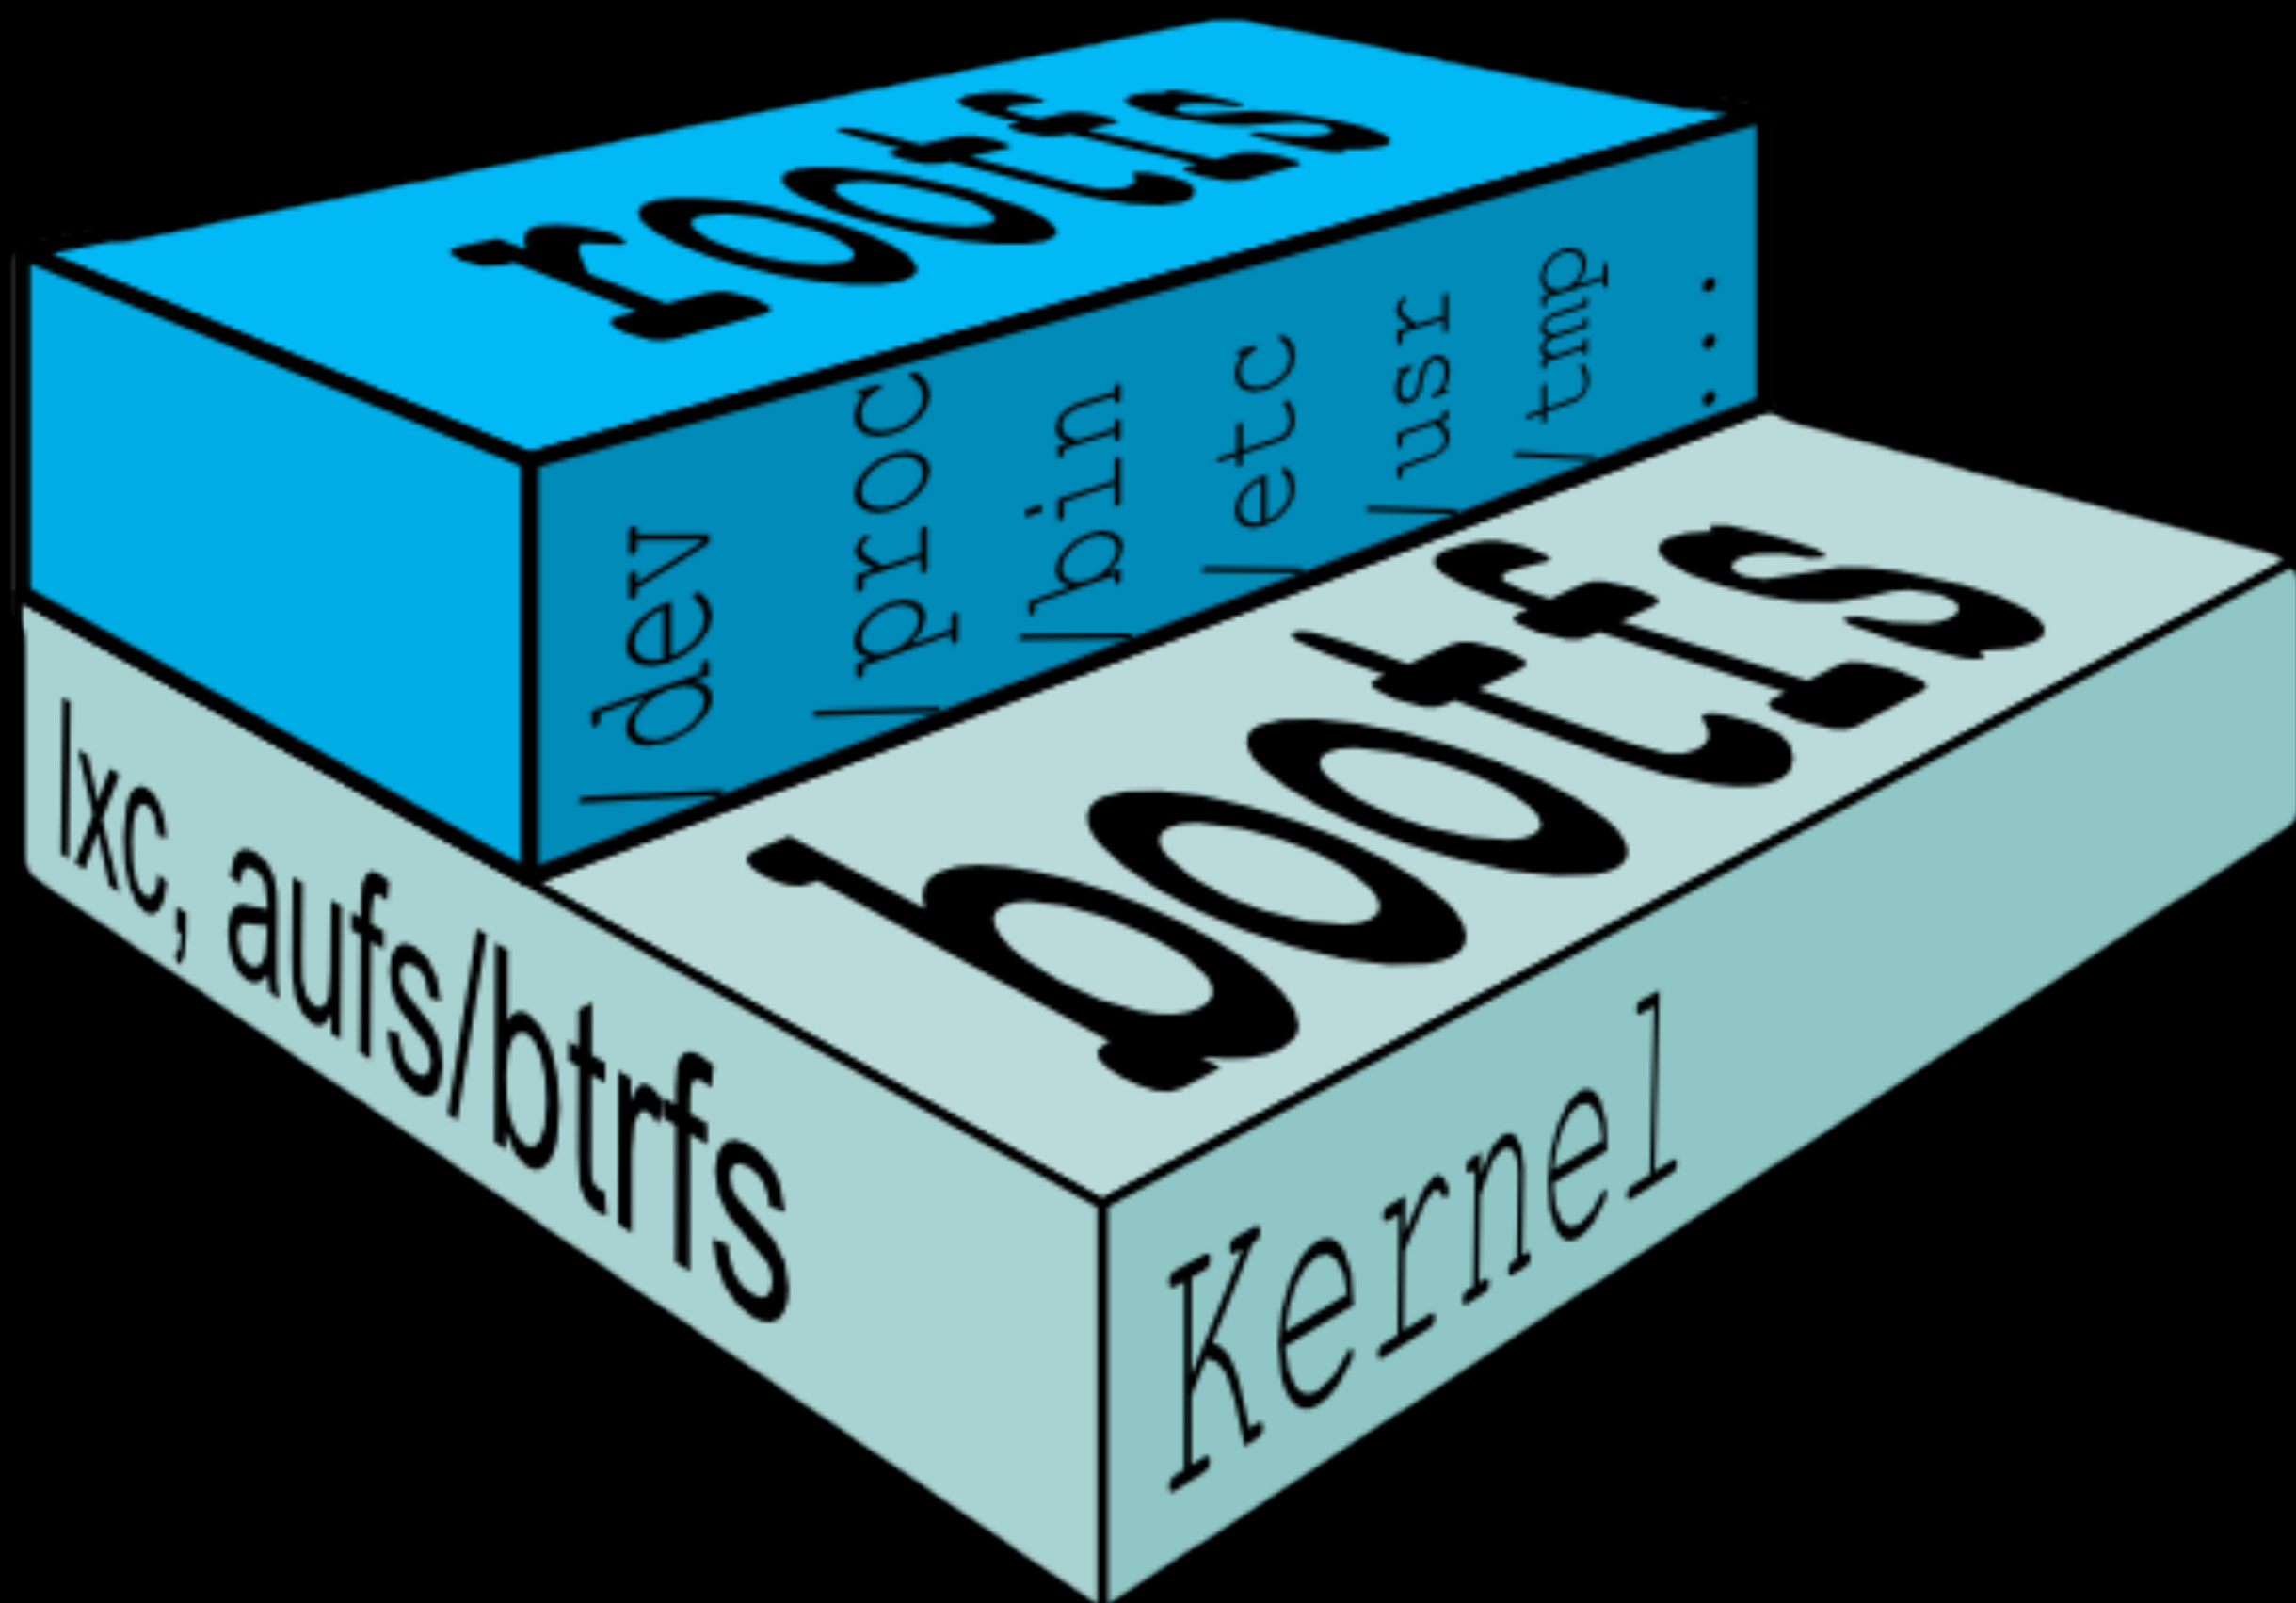
- Runs on most Linux distros
- Boot2Docker for OSX and windows
- Uses Linux Kernel features
- Storage is provided by Union File systems
- Container format

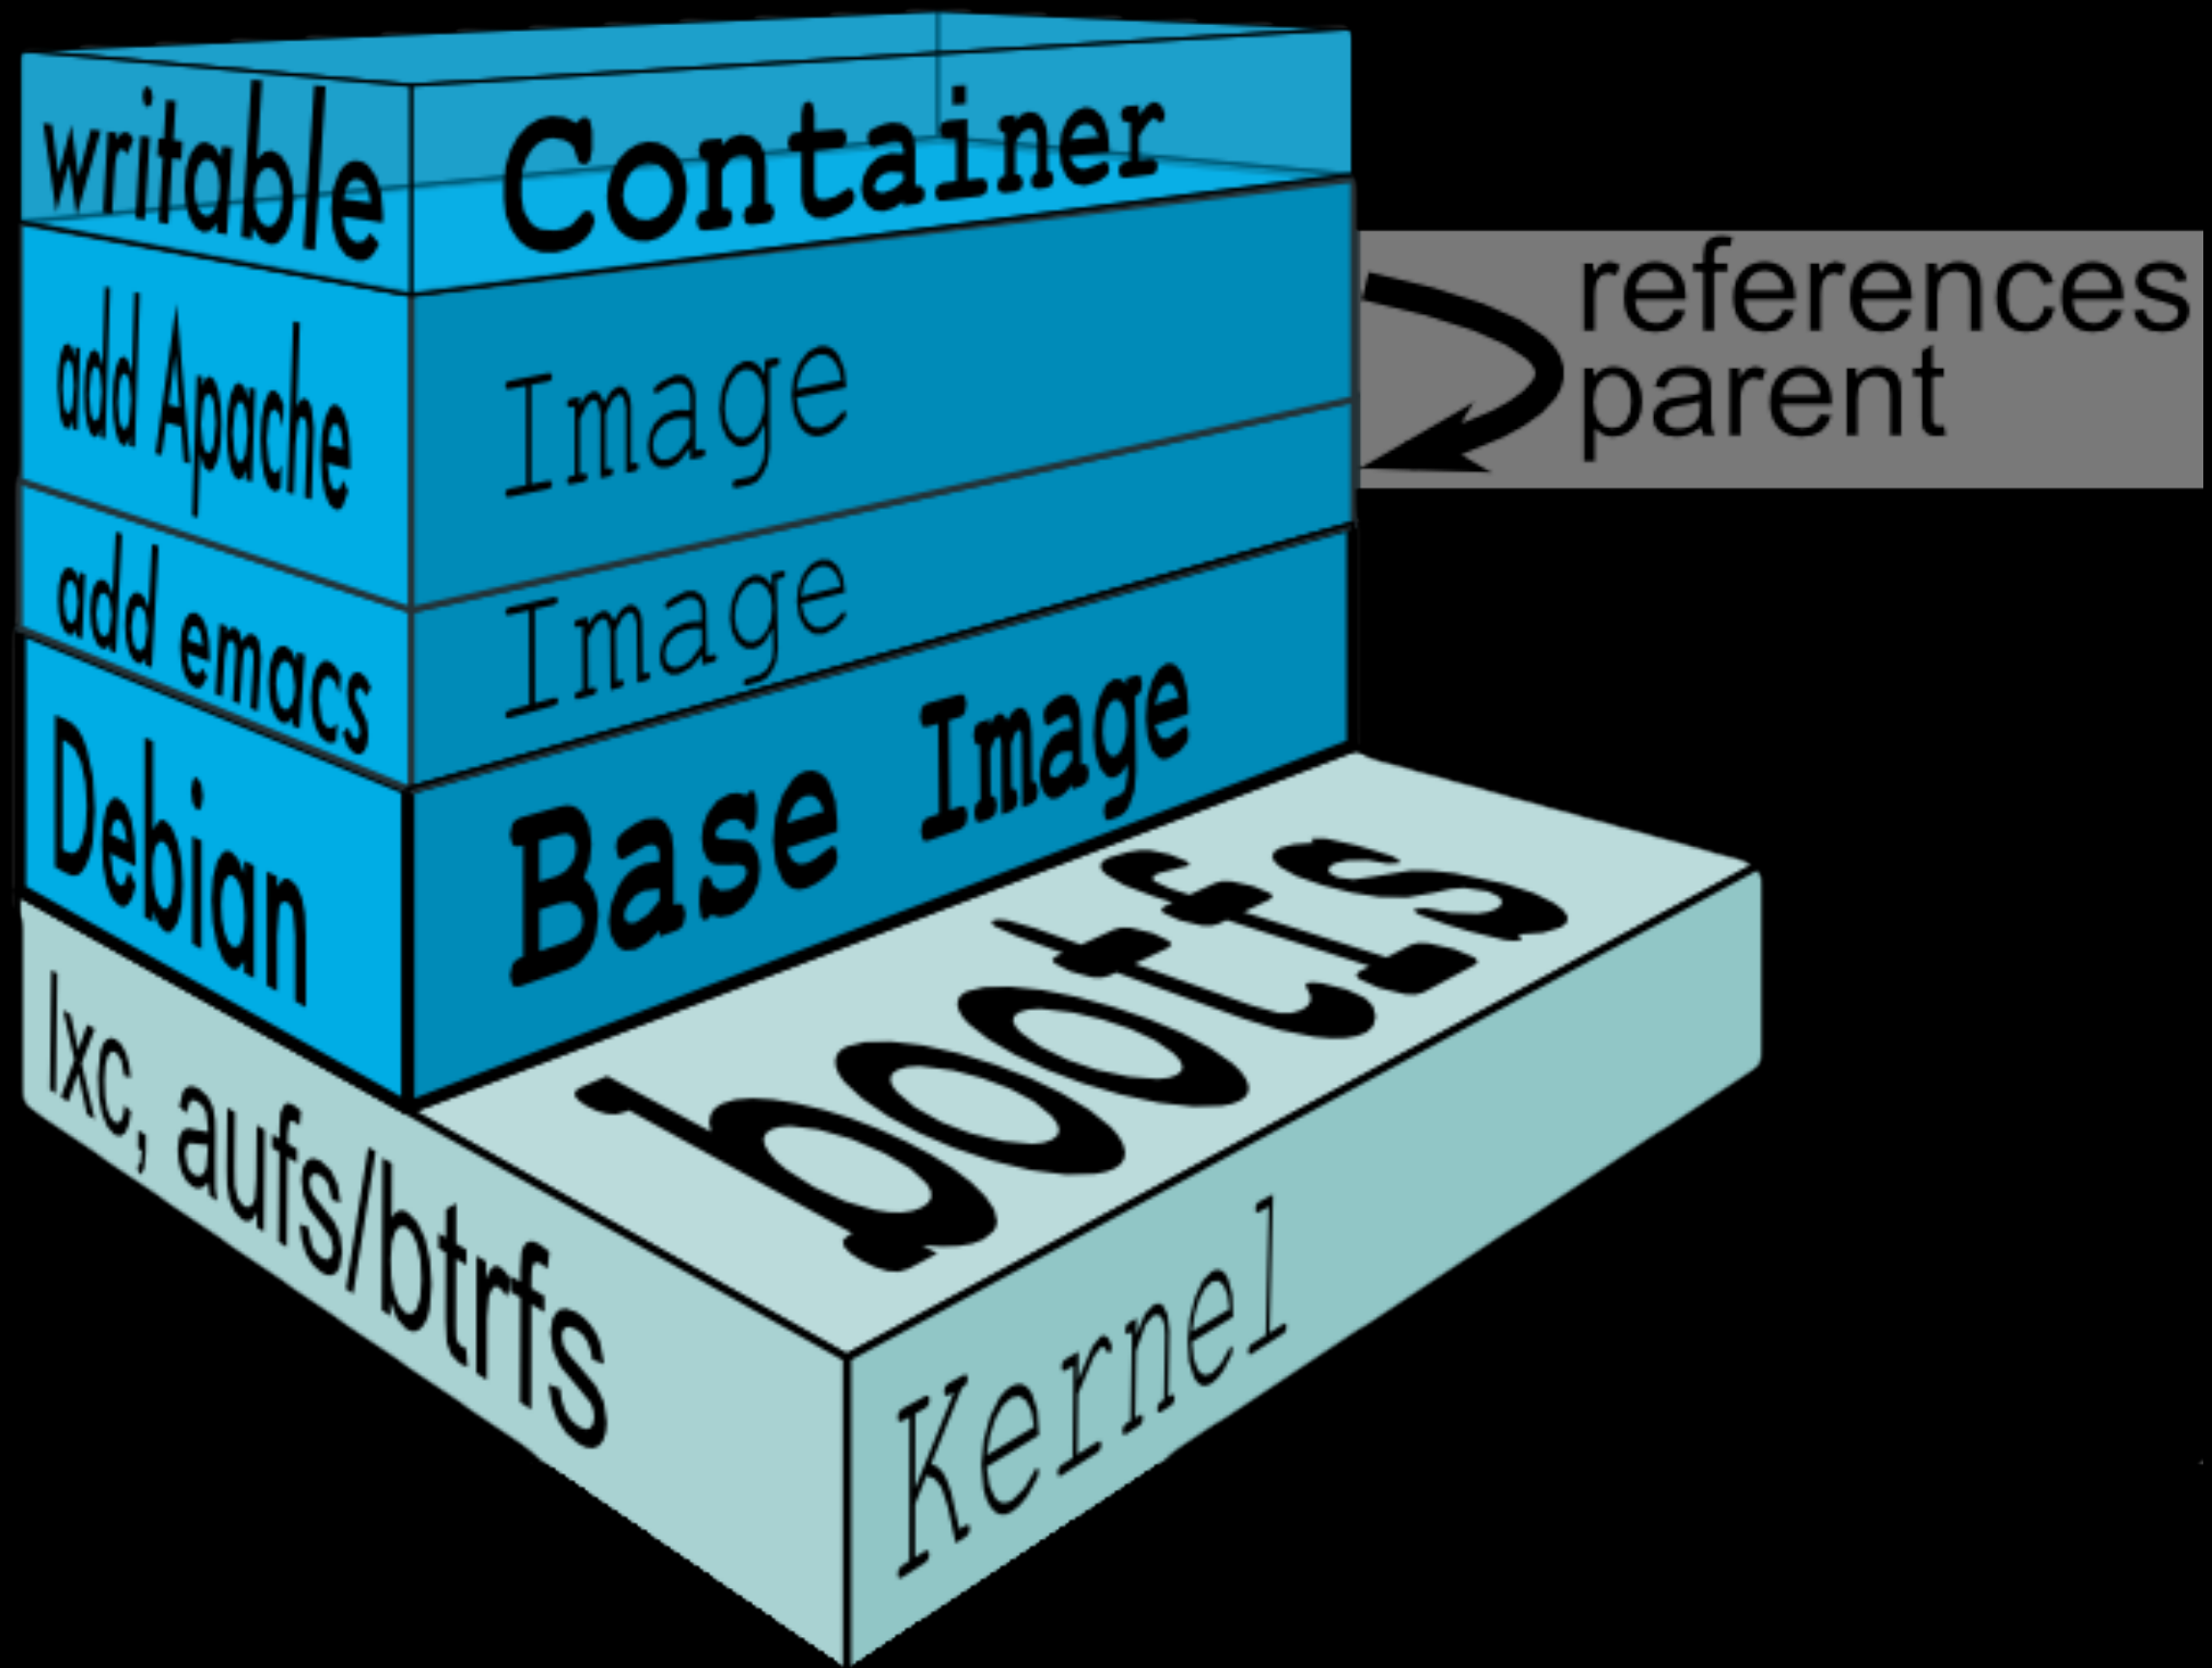
Docker Architecture

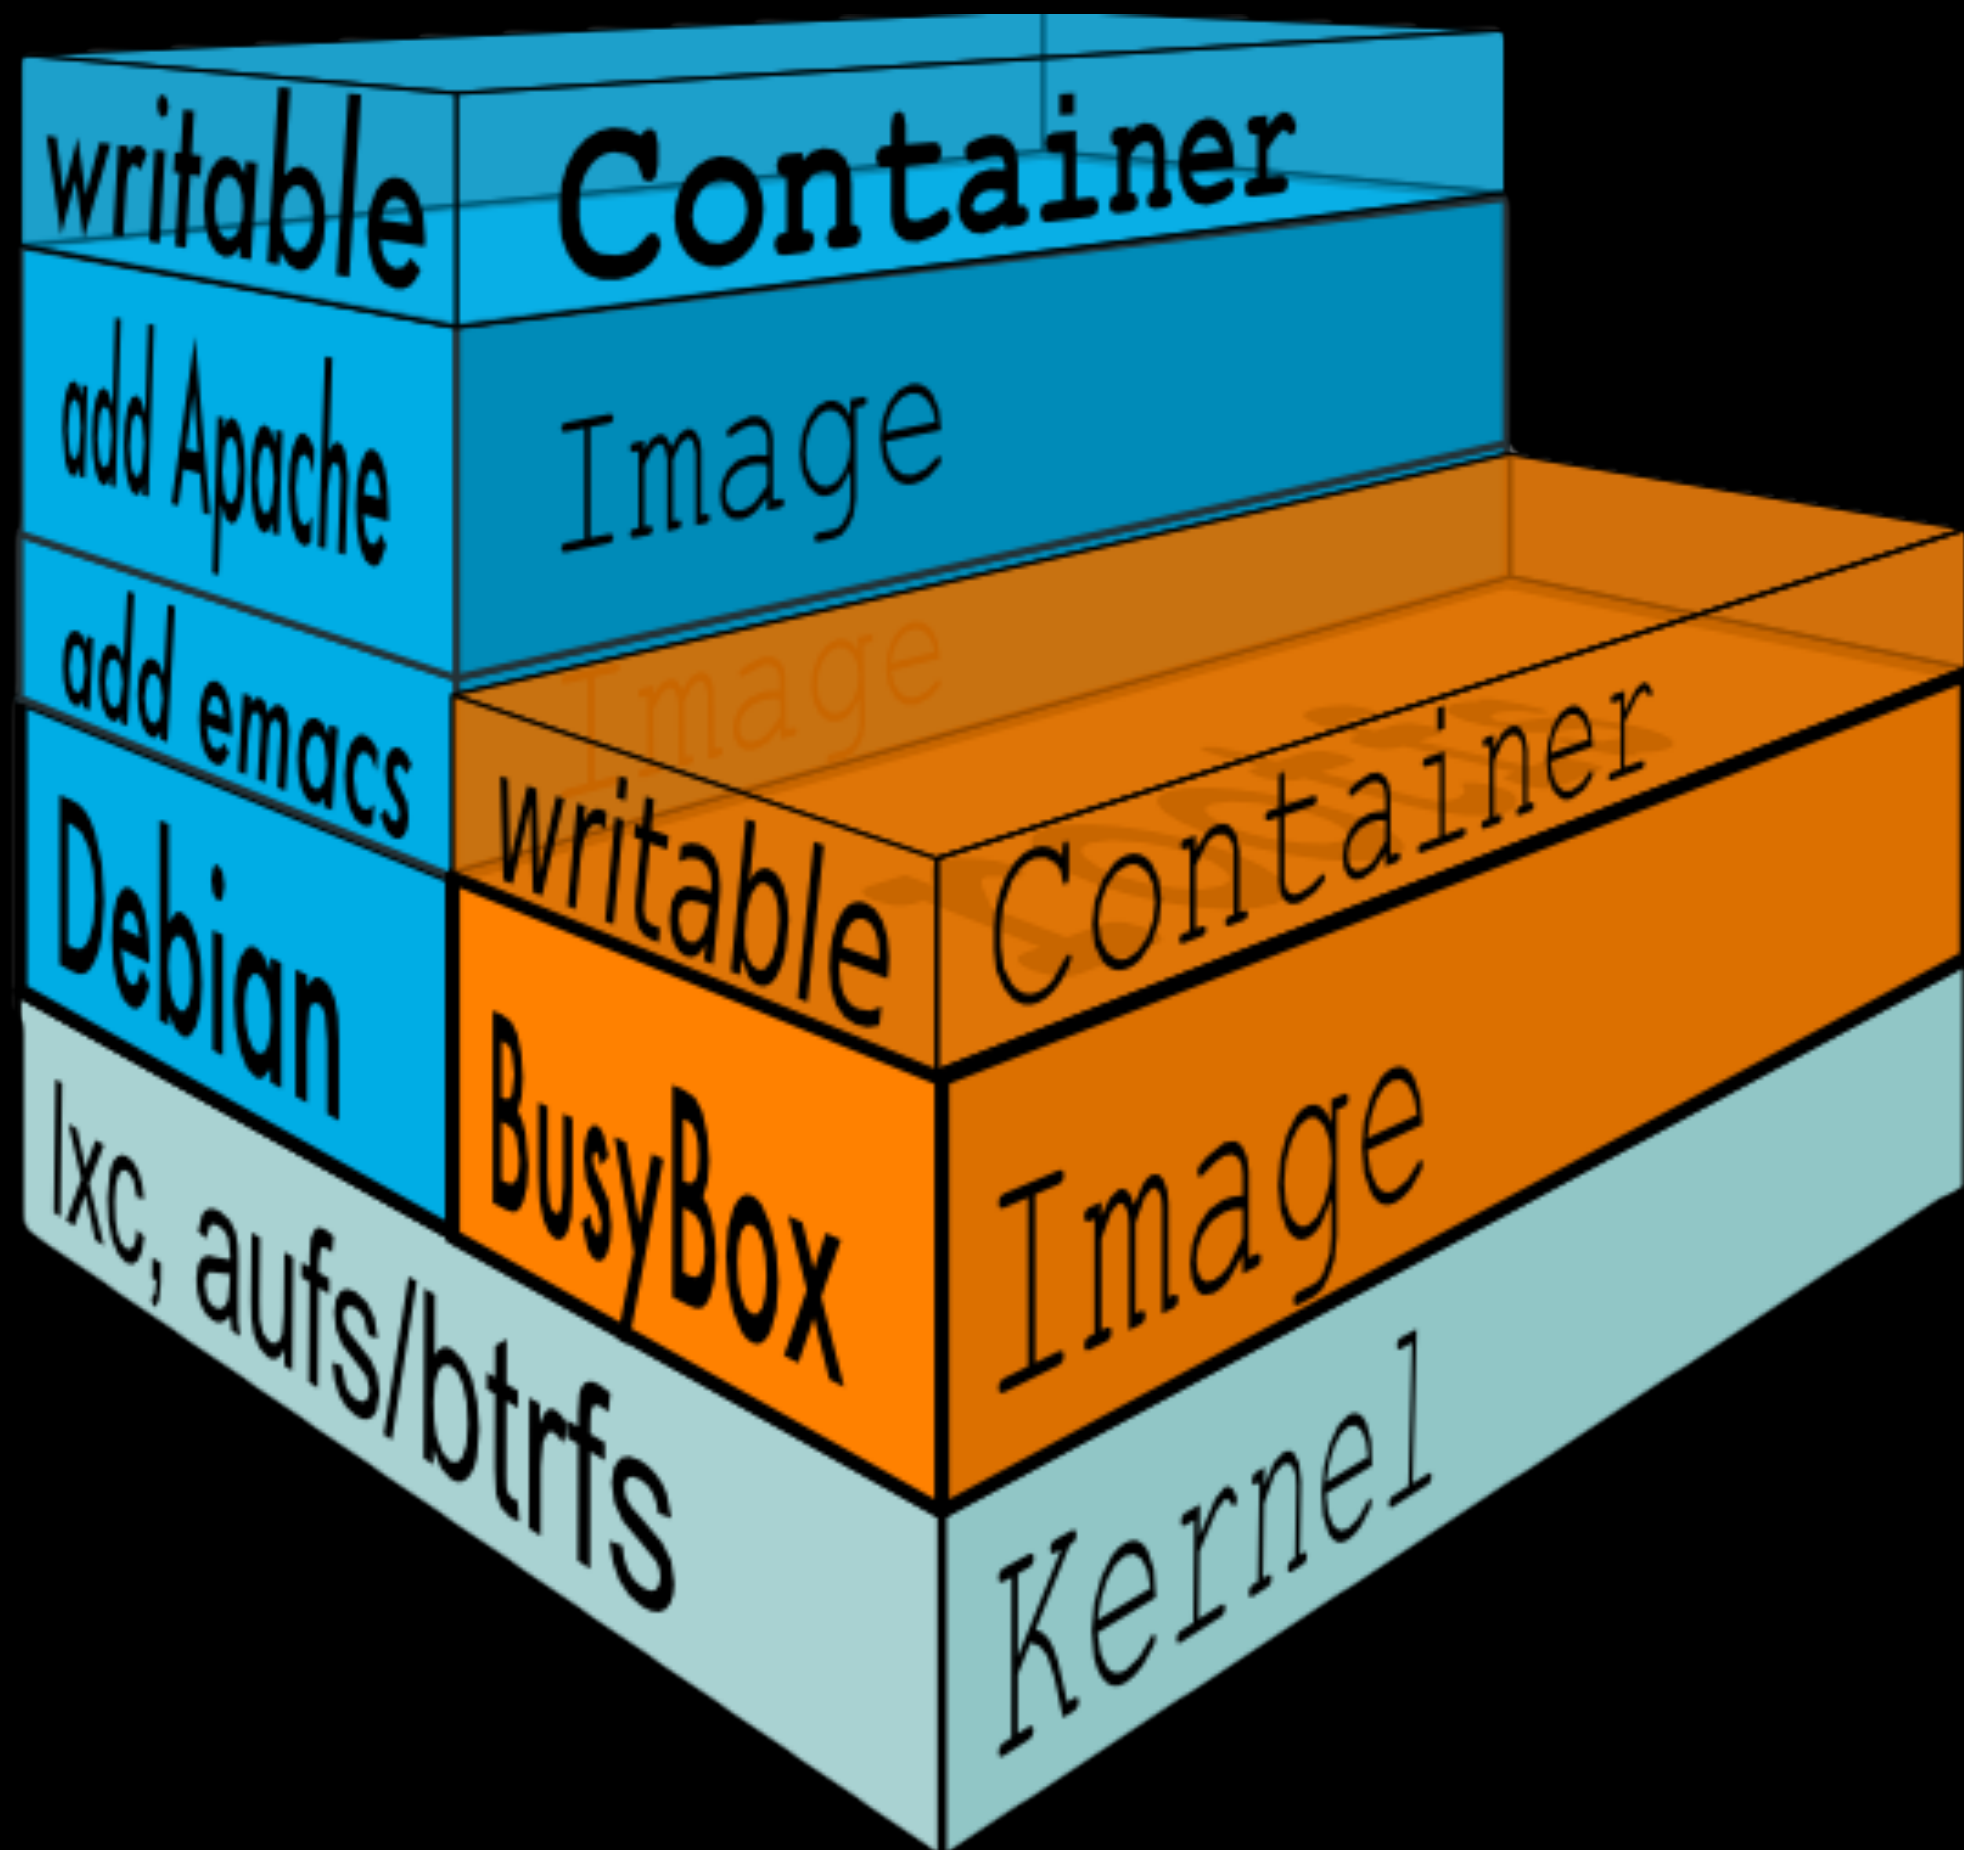


Docker Basics

- Image
- Container
- Registry (like <http://index.docker.io>)







Images

- Images are the source of containers
- Every container is launched from an image
- Images are "built" and layered
- Images are stored in a registry

Layers

- Each image is made up of layers
- Like Linux the bottom layer is a root file system
- Docker uses union file system mounting
- Top layer is writeable and created when a container is launched

Containers

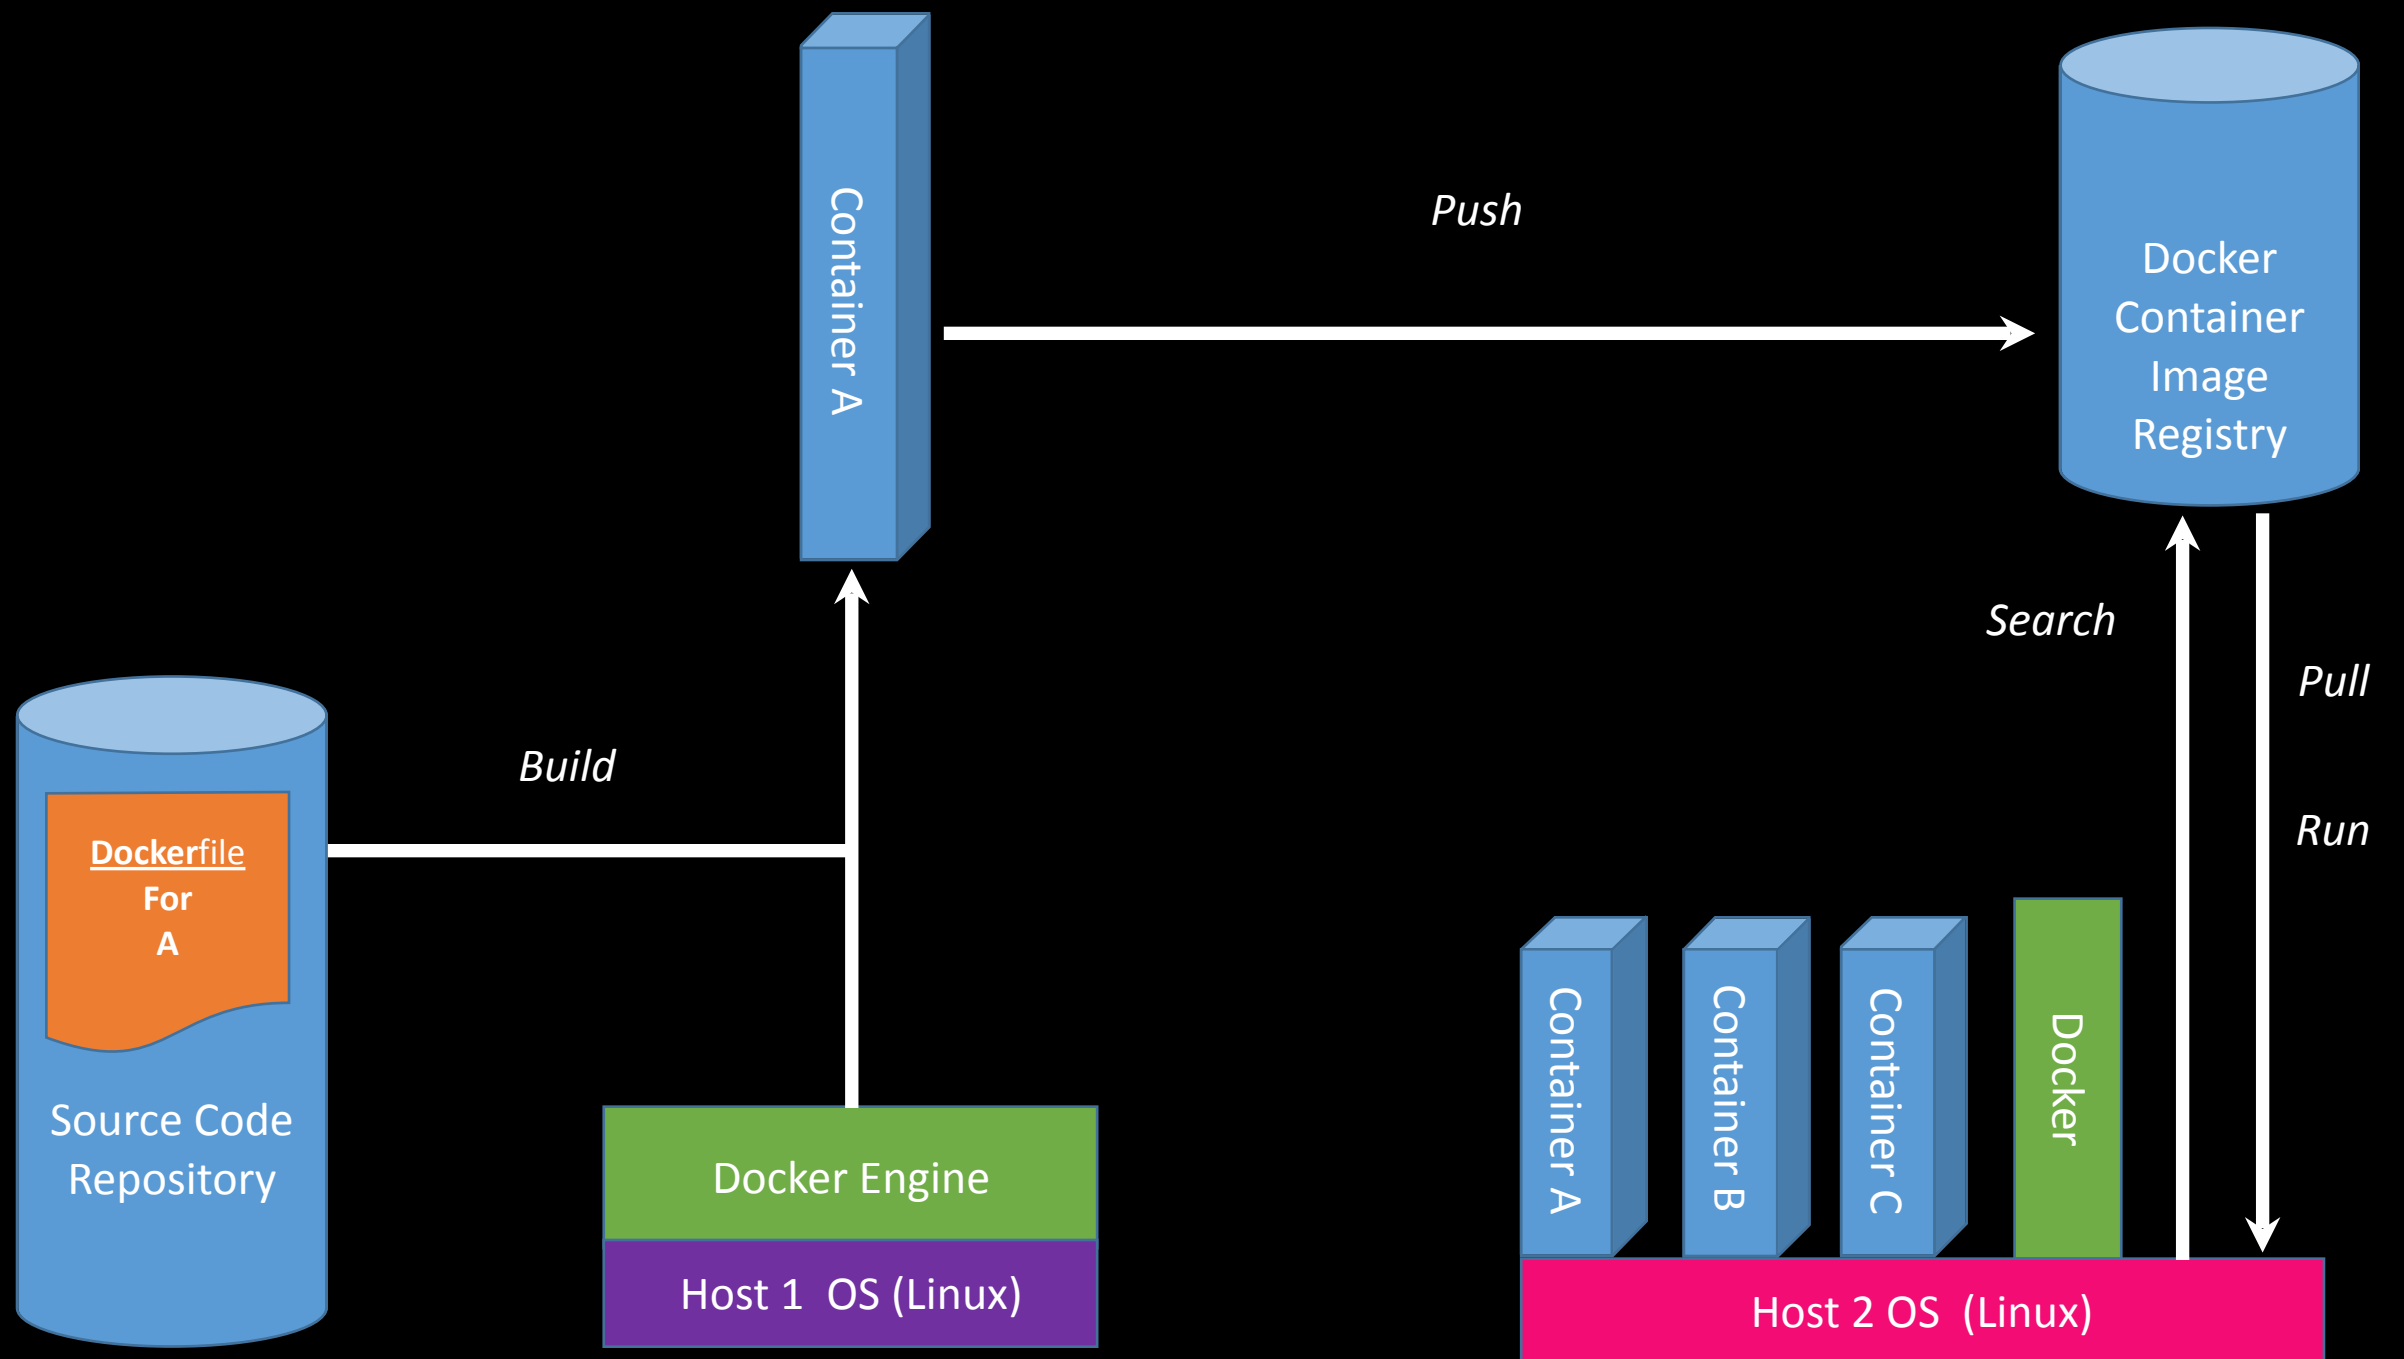
- Containers are launched from images
- They contain one or more running processes
- Can be started, stopped, restarted and killed

Images vs. Containers

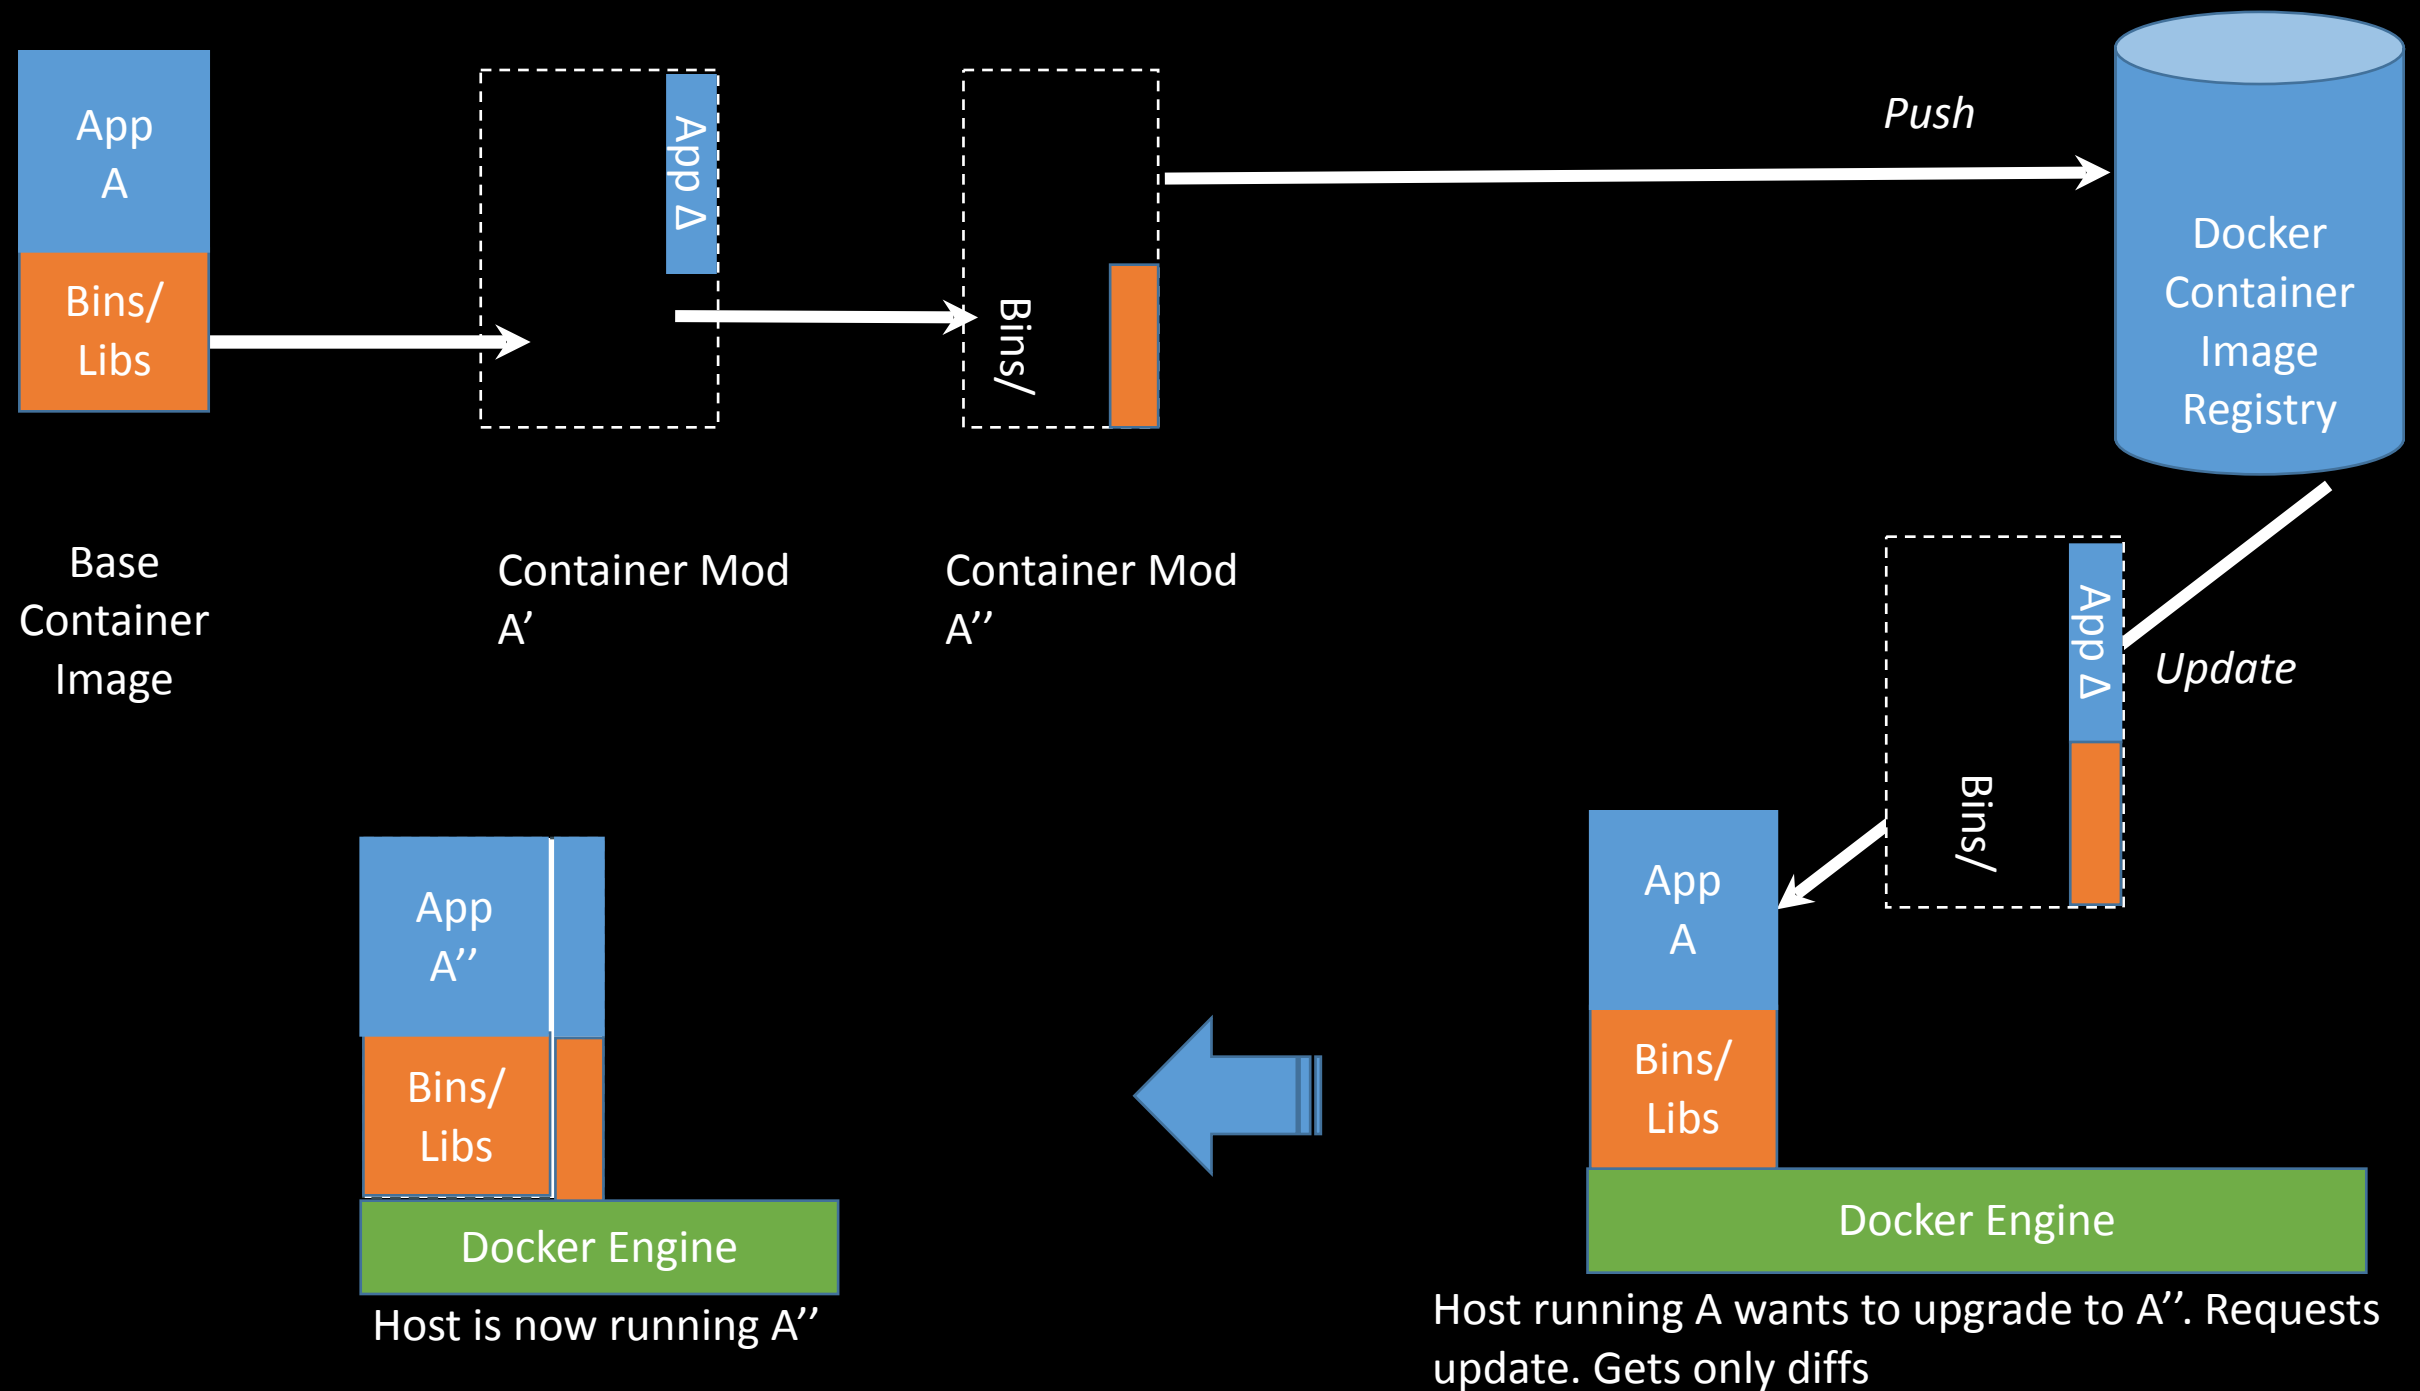
- An image is a stopped container
- You make an image by committing a container
- Images are readonly, containers are writable
- Images are building, containers are running

Build (**Image**)
Ship (by **Registry**)
Run (**Container**)

Workflow



Changes and Updates



Commands

- \$ **docker search** ubuntu
- \$ **docker pull** <user>/<repo>
- \$ **docker run** ubuntu /bin/echo hello world
- \$ id=\$(docker ps -q -l)
- \$ **docker stop** \$id
- \$ **docker start** \$id
- \$ **docker commit** \$CONTAINER_ID <user>/<repo>
- \$ **docker push** <user>/<repo>

Dockfile

Automating Builds

```
FROM ubuntu
MAINTAINER YOUR_NAME YOUR_EMAIL
RUN echo "deb http://archive.ubuntu.com/ubuntu precise main universe" > /etc/apt/sources.list
RUN apt-get update
RUN apt-get install -y memcached
ENTRYPOINT ["memcached"]
USER daemon
EXPOSE 11211
```

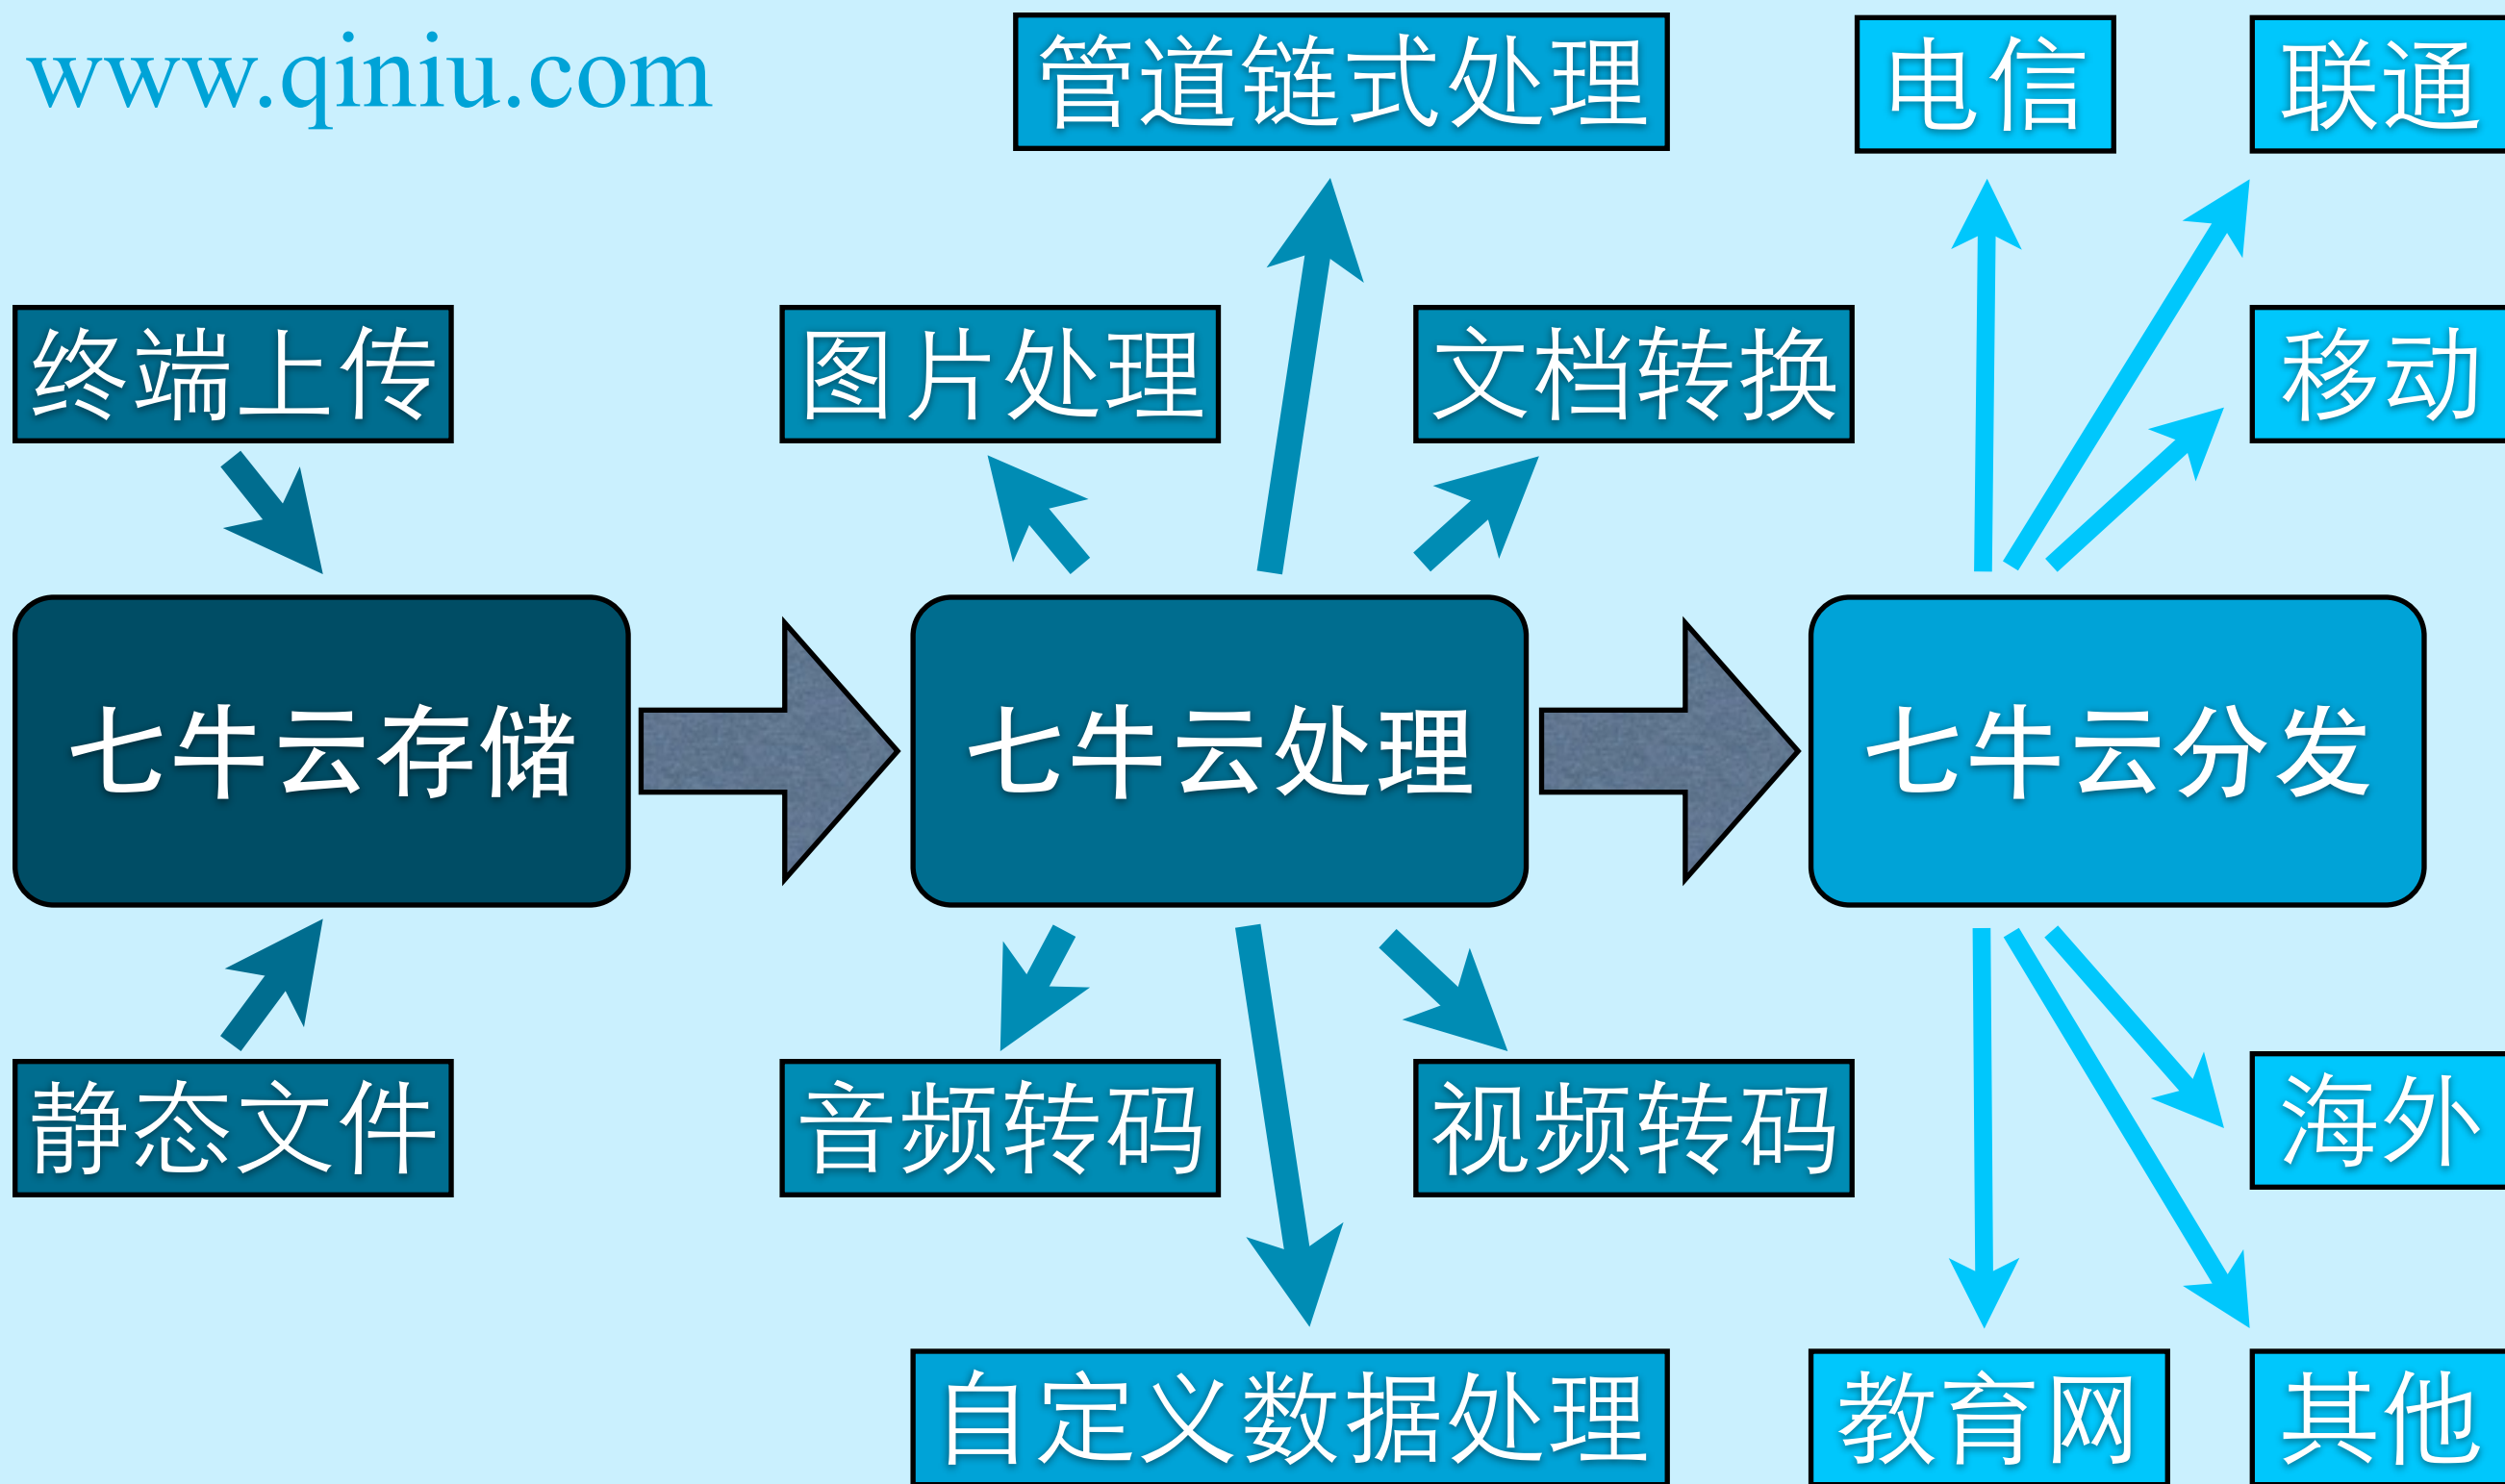
```
$ docker build -t memcached - < Dockerfile
$ docker push <user>/<repo>
```


Resources

- <https://www.docker.io/>
- <https://github.com/dotcloud/docker>
- <http://bit.ly/dockersources>
- <http://lwn.net/Articles/199643/>
- <http://lwn.net/Articles/236038/>
- http://en.wikipedia.org/wiki/Operating_system-level_virtualization
- <https://linuxcontainers.org/>
- <http://en.wikipedia.org/wiki/Cgroups>
- <http://en.wikipedia.org/wiki/Aufs>

Qiniu meet Docker

- Qiniu Cloud Storage service full-stack Go (since 2011)
- Docker is also written in Go (since 2011)
- Both Gopher
- Docker's containers is user-level, not good for disk use
- Qiniu Cloud Storage is os-level, disk-based
- How could they can be together?



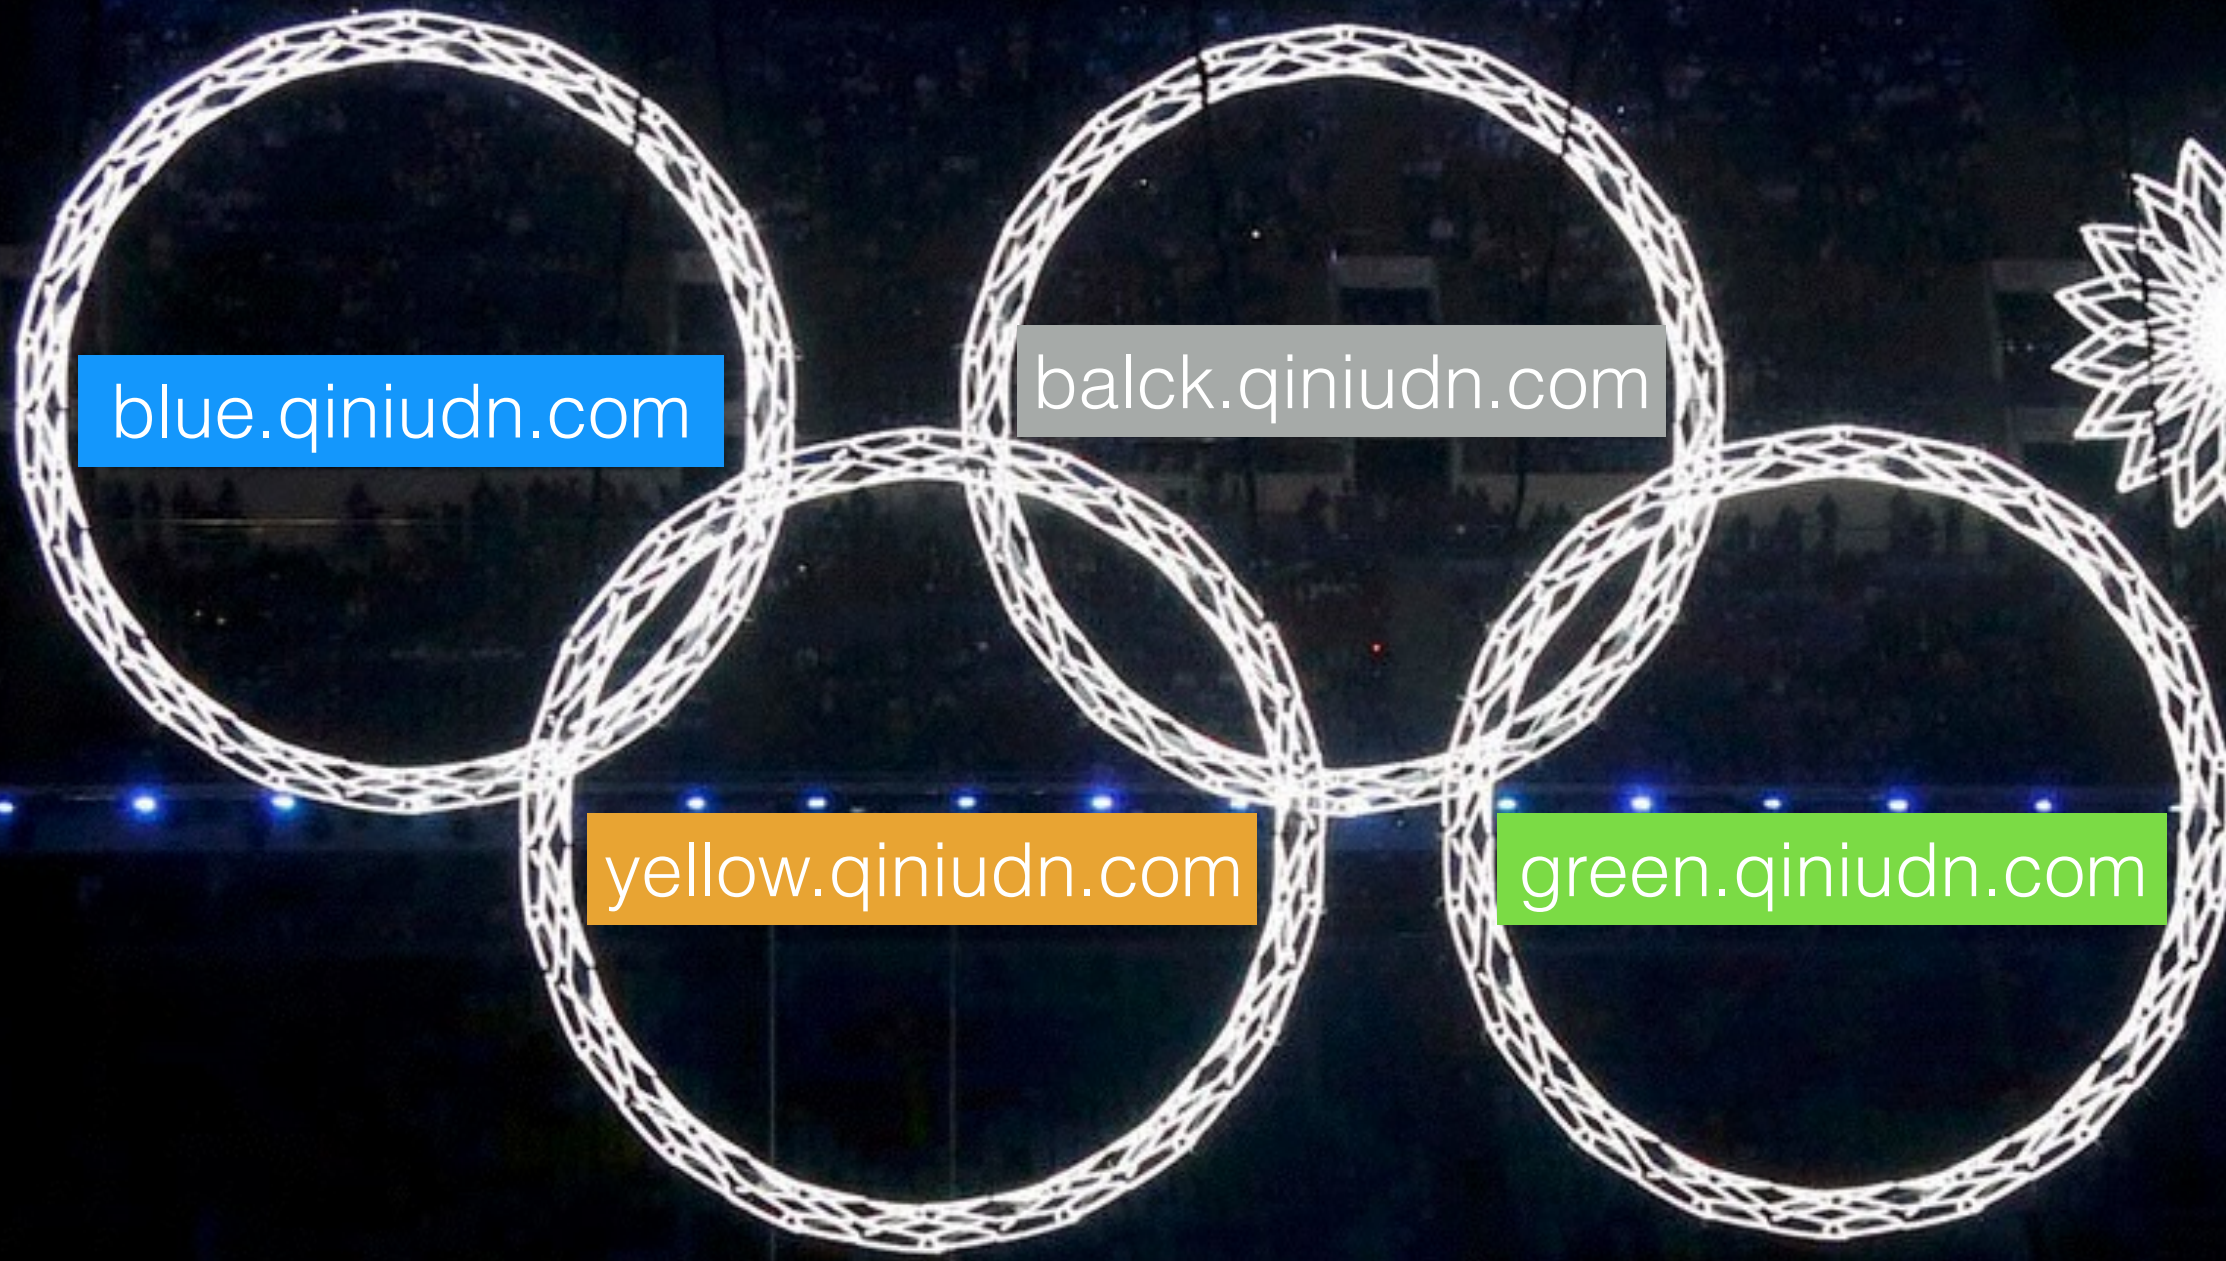
第三方云主机

or

第三方云引擎

or

自有服务器



blue.qiniudn.com

balck.qiniudn.com

yellow.qiniudn.com

green.qiniudn.com

Q & A

- Twitter: why404
- Wechat: why404

THANK YOU !