

Advanced Python Objects

Reuven M. Lerner, PhD
reuven@lerner.co.il

__add__

- Add two elements together
- __add__ gets two parameters
 - self
 - Another object
- How do you add them? That's up to you!
- Typically, you'll return a new instance of the same class

Example

```
class Foo(object):  
    def __init__(self, x):  
        self.x = x  
  
    def __add__(self, other):  
        return Foo(self.x + other.x)  
  
    def __repr__(self):  
        return "Instance of f, x = {}".format(self.x)  
  
f1 = Foo(10)  
f2 = Foo(20)  
print(f1 + f2)
```

`__iadd__`

- We can also address what happens when `+=` is invoked
- This is not the same as `__add__`!
- The expectation is that you'll change the state of the current object, and then return `self`.

Example

```
class Foo(object):  
    def __init__(self, x):  
        self.x = x  
    def __iadd__(self, other):  
        self.x += other.x  
        return self  
    def __repr__(self):  
        return "Instance of f, x = {}".format(self.x)  
  
f1 = Foo(10)  
f2 = Foo(20)  
f1 += f2  
print("Now f1 = {}".format(f1))
```

Making `__add__` flexible

- Remember, the second argument can be anything
- If we want, we can play games with that — checking for attributes

Example

```
class Foo(object):  
    def __init__(self, x):  
        self.x = x  
  
    def __add__(self, other):  
        if hasattr(other, 'x'):  
            return Foo(self.x + other.x)  
        else:  
            return Foo(self.x + other)  
  
    def __repr__(self):  
        return "Instance of f, x = {}".format(self.x)  
  
f1 = Foo(10)  
f2 = Foo(20)  
  
print("f1 + f2 = {}".format(f1 + f2))  
print("f1 + 10 = {}".format(f1 + 50))
```

Reversible?

- What if we now say

```
print("10 + f1 = {}".format(50 + f1))
```

- What will Python do?


```
f1 + 10 = Instance of f, x = 60
```

```
Traceback (most recent call last):
```

```
File "./foo.py", line 29, in <module>
```

```
    print("10 + f1 = {}".format(50 + f1))
```

```
TypeError: unsupported operand type(s) for  
+: 'int' and 'Foo'
```

`__radd__`

- Auto-reversing: We get self and other, and now invoke our previously defined `__add__`:

```
def __radd__(self, other):  
    return self.__add__(other)
```

How does this work?

- Python tries to do it the usual way
- It gets NotImplemented back
 - Not an exception!
 - An instance of NotImplementedType!
- Python then tries (in desperation) to turn things around... and thus, `__radd__`

Type conversions

- You probably know about `__str__` already
- But what about `__int__`? Or even `__hex__`?

Example

```
def __int__(self):  
    return int(self.x)  
  
def __hex__(self):  
    return hex(int(self.x))
```

Boolean conversions

- Your object will always be considered True in a boolean context, unless you define `__nonzero__` (`__bool__` in Python 3)

Example

```
class Foo(object):
```

```
    pass
```

```
f = Foo()
```

```
>>> bool(f)
```

```
True
```

But with `__nonzero__`...

```
class Foo(object):  
    def __init__(self, x):  
        self.x = x  
    def __nonzero__(self):  
        return bool(self.x)
```

```
>>> f = Foo(1)
```

```
>>> bool(f)
```

```
True
```

```
>>> f = Foo(0)
```

```
>>> bool(f)
```

```
False
```


Format

- Originally, Python used the % syntax for pseudo-interpolation:

```
name = 'Reuven'
```

```
print('Hello %s' % name)
```

- But there are lots of problems with it, so it's better to use str.format

str.format

```
name = 'Reuven'
```

```
print('Hello {0}'.format(name))
```

Pad it!

```
name = 'Reuven'
```

```
print('Hello, {0:20}!'.format(name))
```

Move right

```
print('Hello, {0:>20}!'.format(name))
```

Custom formats

- It turns out that our objects can also handle these custom formats!
- If we define `__format__` on our object, then it'll get the format code (i.e., whatever comes after the :)
- We can then decide what to do with it

```
class Person(object):

    def __init__(self, given, family):

        self.given = given

        self.family = family

    def __format__(self, format):

        if format == 'familyfirst':

            return "{} {}".format(self.family, self.given)

        elif format == 'givenfirst':

            return "{} {}".format(self.given, self.family)

        else:

            return "BAD FORMAT CODE"
```

Using it

```
>>> p = Person('Reuven', 'Lerner')
```

```
>>> "Hello, {}".format(p)
```

```
'Hello, BAD FORMAT CODE'
```

```
>>> "Hello, {:familyfirst}".format(p)
```

```
'Hello, Lerner Reuven'
```

```
>>> "Hello, {:givenfirst}".format(p)
```

```
'Hello, Reuven Lerner'
```

Pickle

- Pickle allows you to serialize, or marshall objects
- You can then store them to disk, or send them on the network
- Pickle works with most built-in Python data types, and classes built on those types

Pickling simple data

```
import pickle  
  
d = {'a':1, 'b':2}  
  
p = pickle.dumps(d)  
  
new_d = pickle.loads(p)
```

Custom pickling

- Define some magic methods (of course) to customize your pickling:
- `__getstate__` returns a dictionary that should reflect the object. Want to add to (or remove from) what is being pickled? Just modify a copy of `self.__dict__` and return it!
- Don't modify the dictionary itself...

Example

```
class Foo(object):  
    def __init__(self, x):  
        self.x = x  
    def __getstate__(self):  
        odict = self.__dict__.copy()  
        odict['y'] = self.x * 2  
        return odict  
  
>>> f = Foo(10)  
  
>>> p = pickle.dumps(f)  
  
>>> new_f = pickle.loads(p)  
  
>>> vars(new_f)  
{'x': 10, 'y': 20}
```

Equality

- What makes two object equal?
- By default in Python, they're only equal if their ids are equal
- (Yes, every object has a unique ID number — use the "id" function to find it!)
- You change this by defining `__eq__`!

Changing equality

```
class Foo(object):  
    def __init__(self, x):  
        self.x = x  
  
    def __eq__(self, other):  
        return self.x == other.x
```

```
>>> f1 = Foo(10)
```

```
>>> f2 = Foo(10)
```

```
>>> f1 == f2
```

```
True
```

Hashing

- If two objects are equal, then maybe they should hash to the same value, right?
- We can set the `__hash__` attribute to be a method that returns whatever we want

Playing with `__hash__`

```
class Foo(object):  
    def __init__(self, x):  
        self.x = x  
    def __hash__(self):  
        return hash(self.x)
```

```
>>> f = Foo('a')
```

```
>>> hash(f)
```

```
12416037344
```

```
>>> hash('a')
```

```
12416037344
```

```
f = Foo(1)
```

```
>>> hash(f)
```

```
1
```

Attributes

- Objects in Python depend on having attributes
- Attributes are little namespaces associated with objects
- Each object has its own set of attributes

Listing attributes

- We can list attributes with `dir()` on an object
- This returns a list of strings
- That's nice, but how can we use these strings to understand the attributes?

getattr()

- Built-in function that takes two parameters — an object, and an attribute name

```
f.__class__
```

```
__main__.Foo
```

```
getattr(f, '__class__')
```

```
__main__.Foo
```

setattr()

- Not surprisingly, this built-in function sets an attribute's value

```
getattr(p, 'name')
```

```
'Reuven'
```

```
setattr(p, 'name', 'Waldo')
```

```
getattr(p, 'name')
```

```
'Waldo'
```

```
p.__dict__
```

```
{'name': 'Waldo'}
```

What does this do?

```
def foo(): pass

[ name
  for name in dir(p)
  if type(getattr(p, name)) == type(foo) ]
```

__slots__

- Every instance has `__dict__`, a dictionary of name-value pairs defining its attributes
- This can waste a lot of memory if you're creating many instances with few attributes
- Solution: Define `__slots__` as a class-level attribute

__slots__ example

```
class Person(object):
    __slots__ = ['first', 'last']
    def __init__(self, first, last):
        self.first = first
        self.last = last
p1 = Person('Reuven', 'Lerner')
p1.first
    'Reuven'
p1.last
    'Lerner'
p1.xxx = 'yyy'
    AttributeError
```

Properties

- Properties are "magic" attributes
- You define methods that should execute when the attribute is set or retrieved
- To the user, it appears as though they're just setting or getting an attribute — but really, anything could be happening

property()

- The old style of declaring properties worked as follows:
 - Define getter and setter methods, probably with `_private` names
 - Assign a class attribute to the return value of `property()`, which takes getter and setter methods as arguments


```
class Person(object):
    def __init__(self, first_name, last_name):
        self.first_name = first_name
        self.last_name = last_name
    def get_full_name(self):
        return "{} {}".format(self.first_name, self.last_name)
    def set_full_name(self, full_name):
        self.first_name, self.last_name = full_name.split()
    full_name = property(get_full_name, set_full_name)

p = Person('Reuven', 'Lerner')
print p.full_name          # prints "Reuven Lerner"
p.full_name = 'abc def'
print p.full_name          # prints "abc def"
```

@property

- Today, we can (and should) define properties with a decorator
- The syntax is a bit strange:
 - The getting function is decorated with @property
 - The setter function is decorated with @GETTERNAME.setter

```
#!/usr/bin/env python
class Person(object):
    def __init__(self, first_name, last_name):
        self.first_name = first_name
        self.last_name = last_name

    @property
    def full_name(self):
        return "{} {}".format(self.first_name, self.last_name)

    @full_name.setter
    def full_name(self, full_name):
        self.first_name, self.last_name = full_name.split()

p = Person('Reuven', 'Lerner')
print p.full_name
p.full_name = 'abc def'
print p.full_name
```

When use properties?

- Type checking
- Value checking
- Black-boxing getter/setter functions
- Evolving your API from simple attributes to programmed ones

__getattr__

- For more flexibility on attribute access, you can define __getattr__
- This method is called when the attribute being retrieved does not exist
- If the attribute does exist, __getattr__ is not called

__getattr__ example

```
class Foo(object):
    def __init__(self):
        self.x = 100
    def __getattr__(self, attribute):
        return "{} doesn't exist.".format(attribute)

f = Foo()
print "f.x={}".format(f.x)      #100
print "f.y ={}".format(f.y)     # y doesn't exist
```

`__setattr__`

- Not surprisingly, there is also a `__setattr__` method. (There is also a `__delattr__`)
- However, `__setattr__`, if defined, is called for all attributes, not just those that aren't found.
- If you want to set an attribute within `__setattr__`, you must explicitly call `object.__setattr__`. Otherwise, you get an infinite loop!

```
class Foo(object):
    def __init__(self):
        self.x = 100
    def __setattr__(self, attribute, value):
        print "Setting '{}' to '{}'".format(attribute, value)
        return object.__setattr__(self, attribute, value)

f = Foo()                                # Setting x to 100
print "f.x={}".format(f.x)              # 100

f.x = 200                                # Setting x to 200
print "f.x={}".format(f.x)              # 200
```


__getattribute__

- This is the same idea as `__getattr__`, except that it is invoked unconditionally
- `__getattribute__` is called before Python checks to see if the attribute exists
- If you implement both, then only `__getattribute__` will be invoked.

Descriptors

- You can think of descriptors as working like `__getattr__`, but for a single attribute (rather than for all of them)
- You can also think of properties as simple descriptors
- Descriptors allow you to customize the retrieval, assignment, and deleting of the descriptor attribute
- They also let you reuse this functionality across many classes.

Definitions

- A descriptor is a Python class that implements one or more of the following methods:

`__get__`

`__set__`

`__delete__` (not `__del__` !)

When do we use them?

- Descriptor instances are used for instances of other classes' attributes
- So if you have a class with a particular attribute, you can set it to be an integer, string, list, or dict. Or an instance of any other object... including a descriptor instance.

In other words

- If we have a descriptor class named `MyDescriptor`, we can define a new class that uses it as follows:

```
class MyClass(object):  
  
    a = MyDescriptor("a")  
  
    b = MyDescriptor("b")
```

- Now `a` and `b` are class-level attributes of `MyClass`!

Remember

- Descriptors are defined at the class level
- They are class attributes, not instance attributes

Creating a descriptor

- Define a new class, inheriting (of course) from object
- In that class, define `__init__`, `__get__`, `__set__`, and/or `delete`
- On the class that will use the descriptor, assign a class-level attribute to an instance of the descriptor class

Basic descriptors

```
class Descriptor(object):
    def __init__(self):
        self._name = ''
    def __get__(self, instance, owner):
        print "Getting: {}".format(self._name)
        return self._name
    def __set__(self, instance, name):
        print "Setting: {}".format(name)
        self._name = name.title()
    def __delete__(self, instance):
        print "Deleting: {}".format(self._name)
        del self._name
class Person(object):
    name = Descriptor()
```


Using this code

```
user = Person()  
user.name = 'john smith'  
    Setting: john smith  
user.name  
    Getting: John Smith  
    'John Smith'  
del user.name  
    Deleting: John Smith
```

Descriptor example #1

- Let's say we want to have `Person.fullname` (i.e., the user's first and last names)
- We could define a method, as we've done before
- But an attribute is more natural
- Solution: Define `fullname` as a descriptor

Descriptor class

```
class FullName(object):  
    def __get__(self, instance, owner=None):  
        return "{} {}".format(instance.first_name,  
                                instance.last_name)
```

Things to notice

- `__get__` not only gets its own instance (self), but the instance of the object on which it was invoked, and the class of that instance
- If all we want to do is compute a return value, we can merely implement `__get__`
- We're taking advantage of attribute scoping rules here

Why do this?

- In the case of FullName, we could just implement a method, right?
- Yes, definitely.
- But if you find yourself implementing the same things on many classes, a descriptor can help to DRY up your code

Descriptor example #2

- We can use descriptors to cache data
- Now we will use `__init__` to set up our cache
- The first time it is used, our code will compute the result
- Subsequent calls will retrieve the cached data

```

class FullName(object):
    def __init__(self):
        self.cache = ''
    def __get__(self, instance, owner=None):
        if not self.cache:
            print "Caching!"
            self.cache = "{} {}".format(instance.first_name,
                                          instance.last_name)
        return self.cache

class Person(object):
    def __init__(self, first_name, last_name):
        self.first_name = first_name
        self.last_name = last_name
        fullname = FullName()

p1 = Person('Reuven', 'Lerner')
print p1.fullname      # Reuven Lerner
p2 = Person('Atara', 'Lerner-Friedman')
print p2.fullname      # Reuven Lerner -- uh oh!

```

What happened?

- Our caching worked, but only for the first instance!
- That's because the descriptor is instantiated only once per class! Storing the data in `self.cache` will only work once.
- Let's try a slightly different technique...


```

class FullName(object):
    def __init__(self):
        self.cache = { }
    def __get__(self, instance, owner=None):
        if id(instance) not in self.cache:
            print "Caching!"
            self.cache[id(instance)] = "{} {}".format(instance.first_name,
                                                         instance.last_name)
        return self.cache[id(instance)]

class Person(object):
    def __init__(self, first_name, last_name):
        self.first_name = first_name
        self.last_name = last_name
        fullname = FullName()

p1 = Person('Reuven', 'Lerner')
print p1.fullname           # Reuven Lerner

p2 = Person('Atara', 'Lerner-Friedman')
print p2.fullname           # Atara Lerner-Friedman

```

Multiple inheritance

- Python lets you inherit from multiple bases
- Just specify them in the class definition
- You then have access to methods defined in the base classes

```
class AccountHolder(object):
    def __init__(self, id_number, balance):
        self.id_number = id_number
        self.balance = balance
    def current_balance(self):
        return self.balance
class Person(object):
    def __init__(self, first_name, last_name):
        self.first_name = first_name
        self.last_name = last_name
    def fullname(self):
        return "{} {}".format(self.first_name,
                                self.last_name)
class CEO(Person, AccountHolder):
    def __init__(self, first_name,
                  last_name, id_number, balance):
        Person.__init__(self, first_name, last_name)
        AccountHolder.__init__(self, id_number, balance)
```

Notice `__init__`

- Kind of clunky, right?
- That's because we need to initialize our class with each of its bases
- One way to avoid some (but not all) of this mess is to use `**kwargs` as arguments — then we can pass a dict to all versions of `__init__`

Using our class

```
c = CEO('Reuven', 'Lerner', 123, 1000)
print "Current balance: {}".format(c.current_balance())
print "Full name: {}".format(c.fullname())
```

What about conflicts?

- What if multiple base classes define the same method name?
- Python uses a depth-first search to resolve the methods as best as possible
- As a general rule, earlier bases go first
- If you have a twisted inheritance scheme, consider using less inheritance

Mixins

- One use of multiple inheritance is mixing
- A mixin is a class designed to be used in multiple inheritance, and whose methods take precedence over other base classes' methods
- This lets you replace one or more methods of a class

Example

```
import locale
class CommafyMixin(object):
    def current_balance(self):
        locale.setlocale(locale.LC_ALL, 'en_US')
        return locale.format("%d", self.balance,
                               grouping=True)
```


Using our mixin

- Now, for any class that has a `current_balance()` method, we can mix in our `Commafy` class:

```
class CEO(CommafyMixin, Person, AccountHolder):  
    def __init__(self, first_name, last_name, id_number, balance):  
        Person.__init__(self, first_name, last_name)  
        AccountHolder.__init__(self, id_number, balance)
```

When to use mixing

- Only for methods, not for data!
- Provide many options features for a class via mixing
- Provide one feature in many different classes

type()

- Remember type?

```
a = 'abc'
```

```
type('') == type(a)
```

```
True
```

type() also creates types!

```
X = type('X', (), {'cube':cube,  
                  'square':square, 'y':5})
```

```
x = X()
```

```
x.cube(5)  
125
```

```
x.square(5)  
25
```

```
x.y  
5
```

So what?

- This is interesting, but how does it affect me?
- Consider that every class is an object, and every object has a type that determines its template
- Maybe we want to create a different template for some of our objects?

Metaclasses

- Each class has a metaclass
- The metaclass determines the behavior of the class
- In most cases, you don't really want to be messing with these
- But even so, they're fun and interesting (if very confusing)

Slightly differently:

- Classes
 - are object factories
 - define the behavior of objects
 - define what it means to be an instance of the class

And:

- Metaclasses
 - are class factories
 - define the behavior of classes
 - define what it means to be a class

Default metaclass

- If all classes have a metaclass, then what is the default metaclass?
- That's right — the default metaclass is `type`
- And we learn the metaclass with `type()`

Changing the default

- You can assign a metaclass other than "type" by setting `__metaclass__`
- The value to which you assign `__metaclass__` needs to be a metaclass

Creating a metaclass

- Just subclass "type" instead of "object"

```
class MetaClass(type): pass
```

Now we can use it!

```
class Foo(object):  
    __metaclass__ = MetaClass  
f = Foo()  
  
type(f)  
    <class '__main__.Foo'>  
type(Foo)  
    <class '__main__.MetaClass'>  
type(MetaClass)  
    <type 'type'>
```

Making our metaclass more interesting

- We can define `__init__` on our metaclass
- This is what will be executed when a new class is created!
- That's where metaclasses shine, letting us execute code automatically when a class is *defined*, not just when it is instantiated.

metaclass `__init__`

- It receives four parameters (just as we passed to `type`):
 - `cls` — class
 - `name` — new class name (string)
 - `bases` — tuple of superclasses
 - `dct` — attributes to set

Metaclass definition

```
class CustomMetaclass(type):
    def __init__(cls, name, bases, dct):
        cls.x = 10
        print "Created {} using CustomMetaclass".format(name)
        super(CustomMetaclass, cls).__init__(name,
                                              bases, dct)
```

Using our metaclass

```
class BaseClass(object):  
    __metaclass__ = CustomMetaclass  
class Subclass1(BaseClass):  
    pass
```

Creating class BaseClass using CustomMetaclass

Creating class Subclass1 using CustomMetaclass

Why metaclasses?

- Replace built-in behavior — for example, replace `dump()` with something else
- Set properties/attributes that will work on all classes
- Apply a decorator to some or all attributes on the new class
- Write a DSL (domain-specific language)

Why not metaclasses?

- They take a while to understand — generally considered "black magic"
- Most of what you want to do can probably be done more easily with decorators
- Hard to read and debug code written with them

But why metaclasses?

- They're cool and fun!
- They help you to understand the Python object system
- In some cases, they can reduce code and increase readability

Example: Permissions

```
import os
class PermissionMetaclass(type):
    def __init__(cls, name, bases, dct):
        print "Adding username"
        username = os.environ.get('MC_USERNAME', None)
        cls.username = username
        super(PermissionMetaclass, cls).__init__(name,
                                                    bases,
                                                    dct)
```

Using our metaclass

```
$ export MC_USERNAME=blah
```

```
class Foo(object):  
    __metaclass__ = PermissionMetaclass  
    def __init__(self):  
        self.stuff = 'abc'
```

```
Adding username  
Foo.username  
    'blah'
```

What did we do?

- Every new instance of our metaclass will now have the "username" attribute
- But remember, we only have one chance to do something — at class-creation time!

Playing more

```
import os

class MyAbs(type):
    def __init__(cls, name, bases, dct):
        for key in dct:
            print "Class defines {}".format(key)
        super(MyAbs, cls).__init__(name, bases, dct)
```

Using this metaclass

```
>>> from myabs import *
>>> class Bar(object):
...
...
...
...
...
...
Class defines __module__
Class defines stuff
Class defines __metaclass__
Class defines __init__
```


Subclasses

```
>>> class Bar2(Bar):  
    pass
```

Class defines `__module__`

Enforcement!

```
import os  
class MyAbsError(RuntimeError):  
    pass
```

```

class MyAbs(type):
    def __init__(cls, name, bases, dct):
        required_methods = ['a', 'b', 'c']

        failed_creation = False

        for method_name in required_methods:
            print "Checking for method '{}'".format(method_name)
            if not method_name in dct:
                print "No required method '{}'".format(method_name)
                failed_creation = True
                continue
            the_attribute = getattr(cls, method_name)
            if not callable(the_attribute):
                print "'{}' is not a method".format(method_name)
                failed_creation = True

        if failed_creation:
            raise MyAbsError, "No class for you..."
        super(MyAbs, cls).__init__(name, bases, dct)

```

Using MyAbs

```
print "\n\nCreating class 'Stuff'"
class Stuff(object):
    __metaclass__ = MyAbs
    def a(self): pass
    def b(self): pass
    def c(self): pass
print "\n\nCreating class 'Foo'"
class Foo(object):
    __metaclass__ = MyAbs
    def __init__(self):
self.x = 'x'
```

ABCMeta

- A metaclass that lets you ensure certain methods are defined in a class
- Included in the "abc" ("abstract base class") module, which comes with Python

Using ABCMeta

```
from abc import ABCMeta, abstractmethod
class Person(object):
    __metaclass__ = ABCMeta
    @abstractmethod
    def full_name(self): pass
```

Instantiate Person?

```
>>> p = Person()
```

```
Traceback (most recent call last):
```

```
  File "<stdin>", line 1, in <module>
```

```
TypeError: Can't instantiate abstract class
```

```
Person with abstract methods full_name
```

Instantiate subclass?

```
class Employee(Person):  
    pass
```

```
>>> e = Employee()  
Traceback (most recent call last):  
  File "<stdin>", line 1, in <module>  
TypeError: Can't instantiate abstract class  
Employee with abstract methods full_name
```


But define full_name...

```
class Employee(Person):  
    def full_name(self):  
        return "Reuven Lerner"
```

```
>>> e = Employee()
```

```
>>> e.full_name()  
'Reuven Lerner'
```

So remember:

- When we say:

```
class Name(object):
```

```
    x=5
```

- what is really going on?

What happens?

- The class body is executed as if it were a function
- The names in that function (variables + functions) are put into a dictionary
- We get the class's metaclass M

`Name = M(Name, bases, d)`

Metaclass limitations

- The metaclass doesn't enter the picture until after the class body executes!
- We used that here to our advantage!
- Python 3 provides a metaclass keyword to class as a result

Do we need metaclasses?

- More than 50% of metaclass uses are now obviated by decorators (Alex M.)
- They can be useful for new types of behavior that cut across classes
- But again, your first instinct should be a decorator

`__new__` vs. `__init__`

- Newcomers to Python are surprised to find that `__init__` is an instance method
- Shouldn't a constructor be a class method?
- Yes, and that's why we have `__new__`.

__new__

- __new__ gets a type as its first parameter
- __new__ returns a new instance
- If it doesn't return a new instance, then __init__ isn't ever invoked

`__init__`

- Instance method
- Expects to get the instance in self
- Modifies the attributes of the instance
- Inappropriate for immutable types
 - So for those, use `__new__` instead!

When to use `__new__`

- The big place to do it: When subclassing an immutable type
- Otherwise, think carefully before doing so